Solving the Vehicle Routing Problem with Genetic Algorithms

Áslaug Sóley Bjarnadóttir

Preface

This thesis is the final requirement for obtaining the degree Master of Science in Engineering. The work was carried out at the section of Operations Research at Informatics and Mathematical Modelling, Technical University of Denmark. The duration of the project was from the 10th of September 2003 to the 16th of April 2004. The supervisors were Jesper Larsen and Thomas Stidsen.

First of all, I would like to thank my supervisors for good ideas and suggestions throughout the project.

I would also like to thank Sigurlaug Kristjánsdóttir, Hildur Ólafsdóttir and Þórhallur Ingi Halldórsson for correcting and giving comments on the report. Finally a want to thank my fiance Ingólfur for great support and encouragement.

Odense, April 16th 2004

Áslaug Sóley Bjarnadóttir, s991139

Abstract

In this thesis, Genetic Algorithms are used to solve the Capacitated Vehicle Routing Problem. The problem involves optimising a fleet of vehicles that are to serve a number of customers from a central depot. Each vehicle has limited capacity and each customer has a certain demand. Genetic Algorithms maintain a population of solutions by means of a crossover and mutation operators.

A program is developed, based on a smaller program made by the author and a fellow student in the spring of 2003. Two operators are adopted from that program; Simple Random Crossover and Simple Random Mutation. Additionally, three new crossover operators are developed. They are named Biggest Overlap Crossover, Horizontal Line Crossover and Uniform Crossover. Three Local Search Algorithms are also designed; Simple Random Algorithm, Non Repeating Algorithm and Steepest Improvement Algorithm. Then two supporting operators Repairing Operator and Geographical Merge are made.

Steepest Improvement Algorithm is the most effective one of the Local Search Algorithms. The Simple Random Crossover with Steepest Improvement Algorithm performs best on small problems. The average difference from optimum or best known values is 4,16 $\pm$1,22 %. The Uniform Crossover with Steepest Improvement Crossover provided the best results for large problems, where the average difference was 11.20$\pm$1,79%. The algorithms are called SRC-GA and UC-GA.

A comparison is made of SRC-GA, UC-GA, three Tabu Search heuristics and a new hybrid genetic algorithm, using a number of both small and large problems. SRC-GA and UC-GA are on average 10,52$\pm$5,48% from optimum or best known values and all the other heuristics are within 1%. Thus, the algorithms are not effective enough. However, they have some good qualities, such as speed and simplicity. With that taken into account, they could make a good contribution to further work in the field.

Contents

Chapter 1

Introduction

The agenda of this project is to design an efficient Genetic Algorithm to solve the Vehicle Routing Problem. Many versions of the Vehicle Routing Problem have been described. The Capacitated Vehicle Routing Problem is discussed here and can in a simplified way be described as follows: A fleet of vehicles is to serve a number of customers from a central depot. Each vehicle has limited capacity and each customer has a certain demand. A cost is assigned to each route between every two customers and the objective is to minimize the total cost of travelling to all the customers.

Real life Vehicle Routing Problems are usually so large that exact methods can not be used to solve them. For the past two decades, the emphasis has been on metaheuristics, which are methods used to find good solutions quickly. Genetic Algorithms belong to the group of metaheuristics. Relatively few experiments have been performed using Genetic Algorithms to solve the Vehicle Routing Problem, which makes this approach interesting. Genetic Algorithms are inspired by the Theory of Natural Selection by Charles Darwin. A population of individuals or solutions is maintained by the means of crossover and mutation operators, where crossover simulates reproduction. The quality of each solution is indicated by a fitness value. This value is used to select a solution from the population to reproduce and when solutions are excluded from the population. The average quality of the population gradually improves as new and better solutions are generated and worse solutions are removed.

The project is based on a smaller project developed by the author and Hildur Ólafsdóttir in the course Large-Scale Optimization at DTU in the spring of 2003. In that project a small program was developed, which simulates Genetic Algorithms using very simple crossover and mutation operators. This program forms the basis of the current project.

In this project new operators are designed in order to focus on the geography of the problem, which is relevant to the Capacitated Vehicle Routing Problem. The operators are developed using a trial and error method and experiments are made in order to find out which characteristics play a significant role in a good algorithm. A few Local Search Algorithms are also designed and implemented in order to increase the efficiency. Additionally, an attention is paid to the fitness value and how it influences the performance of the algorithm. The aim of the project is described by the following hypothesis:

It is possible to develop operators for Genetic Algorithms efficient enough to solve large Vehicle Routing Problems.

Problem instances counting more than 100 customers are considered large. What is efficient enough? Most heuristics are measured against the criteria accuracy and speed. Cordeau et al. [4] remark that simplicity and flexibility are also important characteristics of heuristics. The emphasis here is mostly on accuracy. The operators are considered efficient enough if they are able to compete with the best results proposed in the literature. However, an attempt is also made to measure the quality of the operators by the means of the other criteria.

1.1 Outline of the Report

In chapter 2 the theory of the Vehicle Routing Problem and the Genetic Algorithms is discussed. Firstly, the Vehicle Routing Problem is described, the model presented and a review of the literature given among other things. Secondly, the basic concepts of the Genetic Algorithms are explained and different approaches are discussed, e.g. when it comes to choosing a fitness value or a selection method. Then the different types of operators are introduced.

The Local Search Algorithms are presented in chapter 3. Three different algorithms are explained both in words and by a pseudocode. They are compared and the best one chosen for further use.

Chapter 4 describes the development process of the fitness value and the operators. Four crossover operators are explained and in addition; a mutation operator and two supporting operators. All operators are explained both in words and by the means of a pseudocode.

Implementation issues are discussed in chapter 5. This includes information about the computer used for testing, programming language and some relevant methods.

The parameter tuning is described in chapter 6. At first the possible parameters are listed and the procedure of tuning is explained. Then the resulting parameters are illustrated.

Chapter 7 involves the final testing. It starts with a listing of benchmark problems followed by a test description. Then test results are presented. Firstly, different combinations of operators are used to solve a few problems in order to choose the best combination. Secondly, this best combination is applied to a large number of problems. Finally, these results are compared to results presented in the literature.

The results are discussed in chapter 8 and in chapter 9 the conclusion in presented.

1.2   List of Abbreviations

VRP	The Vehicle Routing Problem
GA	Genetic Algorithms
BPP	The Bin Packing Problem
TSP	The Travelling Salesman Problem
SA	Simulated Annealing
DA	Deterministic Annealing
TS	Tabu Search
AS	Ant Systems
NN	Neural Networks
HGA-VRP	A Hybrid Genetic Algorithm
GENI	Generalized Insertion procedure
LSA	Local Search Algorithms
SRA	Simple Random Algorithm
NRA	Non Repeating Algorithm
SIA	Steepest Improvement Algorithm
SRC	Simple Random Crossover
BOC	Biggest Overlap Crossover
GC	First Geography, then Capacity
CG	First Capacity, then Geography
HLC	Horizontal Line Crossover
UC	Uniform Crossover
SRM	Simple Random Mutation
RO	Repairing Operator
GM	Geographical Merge

Chapter 2

Theory

The aim of this chapter is to present the Vehicle Routing Problem (VRP) and Genetic Algorithms (GA) in general. Firstly, VRP is introduced and its model is put forward. Then the nature of the problem is discussed and a review of literature is given. Secondly, GA are introduced and fitness value, selection methods and operators are addressed.

2.1 The Vehicle Routing Problem

2.1.1 The Problem

The Vehicle Routing Problem was first introduced by Dantzig and Ramser in 1959 [12] and it has been widely studied since. It is a complex combinatorial optimisation problem. Fisher [7] describes the problem in a word as to find *the efficient use of a fleet of vehicles that must make a number of stops to pick up and/or deliver passengers or products*. The term *customer* will be used to denote the stops to pick up and/or deliver. Every customer has to be assigned to exactly one vehicle in a specific order. That is done with respect to the capacity and in order to minimise the total cost.

The problem can be considered as a combination of the two well-known optimisation problems; the *Bin Packing Problem* (BPP) and the *Travelling Salesman Problem* (TSP). The BPP is described in the following way: Given a finite set of numbers (the item sizes) and a constant K, specifying the capacity of the bin, what is the minimum number of bins needed?[6] Naturally, all items have to be inside exactly one bin and the total capacity of items in each bin has to be within the capacity limits of the bin. This is known as the best packing version of BPP. The TSP is about a travelling salesman who wants to visit a number of cities. He has to visit each city exactly once, starting and ending in his home town. The problem is to find the shortest tour through all cities. Relating this to the VRP, customers can be assigned to vehicles by solving BPP and the order in which they are visited can be found by solving TSP.

Figure 2.1 shows a solution to a VRP as a graph.

Figure 2.1: A solution to a Vehicle Routing Problem. Node 0 denotes the depot and nodes $1 - 10$ are the customers.

2.1.2   The Model

The most general version of VRP is the *Capacitated Vehicle Routing Problem*, which will be referred to as just VRP from now on. The model for VRP has the following parameters [7]:

n	is the number of customers,
K	denotes the capacity of each vehicle,
d_i	denotes the demand of customer i (in same units as vehicle capacity) and
c_{ij}	is the cost of travelling from customer i to customer j.

All parameters are considered non-negative integers. A homogeneous fleet of vehicles with a limited capacity K and a central depot, with index 0, makes deliveries to customers, with indices 1 to n. The problem is to determine the exact tour of each vehicle starting and ending at the depot. Each customer must be assigned to exactly one tour, because each customer can only be served by one vehicle. The sum over the demands of the customers in every tour has to be within the limits of the vehicle capacity. The objective is to minimise the total travel cost. That could also be the distance between the nodes or other quantities on which the quality of the solution depends, based on the problem to be solved. Hereafter it will be referred to as a cost.

The mathematical model is defined on a graph (N,A). The node set N corresponds to the set of customers C from 1 to n in addition to the depot number 0. The arc set A consists of possible connections between the nodes. A connection between every two nodes in the graph will be included in A here. Each arc $(i,j) \in A$ has a travel cost c_{ij} associated to it. It is assumed that the cost is symmetric, i.e. $c_{ij} = c_{ji}$, and also that $c_{ii} = 0$. The set of uniform vehicles is V. The vehicles have a capacity K and all customers have a demand d_i. The only decision variable is X_{ij}^v:

$$X_{ij}^v = \begin{cases} 1 & \text{if vehicle } v \text{ drives from node } i \text{ to node } j \\ 0 & \text{otherwise} \end{cases} \qquad (2.1)$$

The objective function of the mathematical model is:

$$min \sum_{v \in V} \sum_{(i,j) \in A} c_{ij} X_{ij}^v \tag{2.2}$$

subject to

$$\sum_{v \in V} \sum_{j \in N} X_{ij}^v = 1 \qquad \forall i \in C \tag{2.3}$$

$$\sum_{i \in C} d_i \sum_{j \in N} X_{ij}^v \leq K \qquad \forall v \in V \tag{2.4}$$

$$\sum_{j \in C} X_{0j}^v = 1 \qquad \forall v \in V \tag{2.5}$$

$$\sum_{i \in N} X_{ik}^v - \sum_{j \in N} X_{kj}^v = 0 \qquad \forall k \in C \quad \text{and} \quad \forall v \in V \tag{2.6}$$

$$X_{ij}^v \in \{0, 1\}, \quad \forall (i, j) \in A \quad \text{and} \quad \forall v \in V \tag{2.7}$$

Equation 2.3 is to make sure that each customer is assigned to exactly one vehicle. Precisely one arc from customer i is chosen, whether or not the arc is to another customer or to the depot. In equation 2.4 the capacity constraints are stated. The sum over the demands of the customers within each vehicle v has to be less than or equal to the capacity of the vehicle. The flow constraints are shown in equations 2.5 and 2.6. Firstly, each vehicle can only leave the depot once. Secondly, the number of vehicles entering every customer k and the depot must be equal to the number of vehicles leaving.

An even simpler version could have a constant number of vehicles but here the number of vehicles can be modified in order to obtain smallest possible cost. However, there is a lower bound on the number of vehicles, which is the smallest number of vehicles that can carry the total demand of the customers, $\lceil \frac{\sum_{i \in C} d_i \sum_{j \in N} X_{ij}^v}{K} \rceil$.

2.1.3   VRP in Real Life

The VRP is of great practical significance in real life. It appears in a large number of practical situations, such as transportation of people and products, delivery service and garbage collection. For instance, such a matter of course as being able to buy milk in a store, arises the use of vehicle routing twice. First the milk is collected from the farms and transported to the dairy and when it has been put into cartons it is delivered to the stores. That is the way with most of the groceries we buy. And the transport is not only made by vehicles but also by plains, trains and ships. VRP is everywhere around!

One can therefore easily imagine that all the problems, which can be considered as VRP, are of great economic importance, particularly to the developed nations. The economic

importance has been a great motivation for both companies and researches to try to find better methods to solve VRP and improve the efficiency of transportation.

2.1.4 Solution Methods and Literature Review

The model above describes a very simple version of VRP. In real life, VRP can have many more complications, such as asymmetric travel costs, multiple depots, heterogeneous vehicles and time windows, associated with each customer. These possible complications make the problem more difficult to solve. They are not considered in this project because the emphasis is rather on Genetic Algorithms.

In section 2.1.1 above, it is explained how VRP can be considered a merge of BPP and TSP. Both BPP and TSP are so-called NP-hard problems [6] and [21], thus VRP is also NP-hard. NP-hard problems are difficult to solve and in fact it means that to date no optimal algorithm has been found, which is able to solve the problem in polynomial time [6]. Finding an optimal solution to a NP-hard problem is usually very time consuming or even impossible. Because of this nature of the problem, it is not realistic to use exact methods to solve large instances of the problem. For small instances of only few customers, the *branch and bound* method has proved to be the best [15]. Most approaches for large instances are based on heuristics. Heuristics are approximation algorithms that aim at finding good feasible solutions quickly. They can be roughly divided into two main classes; *classical heuristics* mostly from between 1960 and 1990 and *metaheuristics* from 1990 [12].

The classical heuristics can be divided into three groups; Construction methods, two-phase methods and improvement methods [13]. Construction methods gradually build a feasible solution by selecting arcs based on minimising cost, like the *Nearest Neighbour* [11] method does. The two-phase method divides the problem into two parts; clustering of customers into feasible routes disregarding their order and route construction. An example of a two-phase method is the *Sweep Algorithm* [12], which will be discussed further in section 4.2.3. The *Local Search Algorithms* [1], explained in chapter 3, belong to the improvement heuristics. They start with a feasible solution and try to improve it by exchanging arcs or nodes within or between the routes. The advantage of the classical heuristics is that they have a polynomial running time, thus using them one is better able to provide good solutions within a reasonable amount of time [4]. On the other hand, they only do a limited search in the solution space and do therefore run the risk of resulting in a local optimum.

Metaheuristics are more effective and specialised than the classical heuristics [5]. They combine more exclusive neighbourhood search, memory structures and recombination of solutions and tend to provide better results, e.g. by allowing deterioration and even infeasible solutions [10]. However, their running time is unknown and they are usually more time consuming than the classical heuristics. Furthermore, they involve many parameters that need to be tuned for each problem before they can be applied.

For the last ten years metaheuristics have been researched considerably, producing some effective solution methods for VRP [4]. At least six metaheuristics have been applied to

VRP; *Simulated Annealing (SA)*, *Deterministic Annealing* (DA), *Tabu Search* (TS), *Ant Systems* (AS), *Neural Networks* (NN) and *Genetic Algorithms* (GA) [10]. The algorithms SA, DA and TS move from one solution to another one in the neighbourhood until a stopping criterion is satisfied. The fourth method, AS, is a constructive mechanism creating several solutions in each iteration based on information from previous generations. NN is a learning method, where a set of weights is gradually adjusted until a satisfactory solution is reached. Finally, GA maintain a population of good solutions that are recombined to produce new solutions.

Compared to best-known methods, SA, DA and AS have not shown competitive results and NN are clearly outperformed [10]. TS has got a lot of attention by researches and so far it has proved to be the most effective approach for solving VRP [4]. Many different TS heuristics have been proposed with unequal success. The general idea of TS and a few variants thereof are discussed below. GA have been researched considerably, but mostly in order to solve TSP and VRP with time windows [2], where each customer has a time window, which the vehicle has to arrive in. Although they have succeeded in solving VRP with time windows, they have not been able to show as good results for the capacitated VRP. In 2003 Berger and Barkaoui presented a new Hybrid Genetic Algorithm (HGA-VRP) to solve the capacitated VRP [2]. It uses two populations of solutions that periodically exchange some number of individuals. The algorithm has shown to be competitive in comparison to the best TS heuristics [2]. In the next two subsections three TS approaches are discussed followed by a further discussion of HGA-VRP.

Tabu Search

As written above, to date Tabu Search has been the best metaheuristic for VRP [4]. The heuristic starts with an initial solution x_1 and in step t it moves from solution x_t to the best solution x_{t+1} in its neighbourhood $N(x_t)$, until a stopping criterion is satisfied. If $f(x_t)$ denotes the cost of solution x_t, $f(x_{t+1})$ does not necessarily have to be less than $f(x_t)$. Therefore, a cycling must be prevented, which is done by declaring some recently examined solutions *tabu* or forbidden and storing them in a tabulist. Usually, the TS methods preserve an attribute of a solution in the tabulist instead of the solution itself to save time and memory. Different TS heuristics have been proposed not all with equal success. For the last decade, some successful TS heuristics have been proposed [12].

The *Taburoute* of Gendreau et al. [9] is an involved heuristic with some innovative features. It defines the neighbourhood of x_t as a set of solutions that can be reached from x_t by removing a customer k from its route r and inserting it into another route s containing one of its nearest neighbours. The method uses Generalised Insertion (GENI) procedure also developed by Gendreau et al. [8]. Reinsertion of k into r is forbidden for the next θ iterations, where θ is a random integer in the interval (5,10) [12]. A diversification strategy is used to penalise frequently moved nodes. The Taburoute produces both feasible and infeasible solutions.

The *Taillard's Algorithm* is one of the most accurate TS heuristics [4]. Like Taburoute

it uses random tabu duration and diversification. However, the neighbourhood is defined by the means of λ-interchange generation mechanism and standard insertion methods are used instead of GENI. The innovative feature of the algorithm is the decomposition of the main problem into subproblems.

The *Adaptive Memory* procedure of Rochat and Taillard is the last TS heuristic that will be discussed here. It is probably one of the most interesting novelties that have emerged within TS heuristics in recent years [12]. An adaptive memory is a pool of solutions, which is dynamically updated during the search process by combining some of the solutions in the pool in order to produce some new good solutions. Therefore, it can be considered a generalisation of the genetic search.

A Hybrid Genetic Algorithm

The *Hybrid Genetic Algorithm* proposed by Berger and Barkaoui is able to solve VRP in almost as effective way as TS [2]. Genetic Algorithms are explained in general in the next section. The algorithm maintains two populations of solutions that exchange a number of solutions at the end of each iteration. New solutions are generated by rather complex operators that have successfully been used to solve the VRP with time windows. When a new best solution has been found the customers are reordered for further improvement. In order to have a constant number of solutions in the populations the worst individuals are removed. For further information about the Hybrid Genetic Algorithm the reader is referred to [2].

2.2   Genetic Algorithms

2.2.1   The Background

The *Theory of Natural Selection* was proposed by the British naturalist Charles Darwin (1809-1882) in 1859 [3]. The theory states that individuals with certain favourable characteristics are more likely to survive and reproduce and consequently pass their characteristics on to their offsprings. Individuals with less favourable characteristics will gradually disappear from the population. In nature, the genetic inheritance is stored in *chromosomes*, made of genes. The characteristics of every organism is controlled by the genes, which are passed on to the offsprings when the organisms mate. Once in a while a mutation causes a change in the chromosomes. Due to natural selection, the population will gradually improve on the average as the number of individuals having the favourable characteristics increases.

The *Genetic Algorithms* (GA) were invented by John Holland and his colleagues in the early 1970s [16], inspired by Darwin's theory. The idea behind GA is to model the natural evolution by using genetic inheritance together with Darwin's theory. In GA, the *population* consists of a set of *solutions* or *individuals* instead of chromosomes. A *crossover* operator plays the role of reproduction and a *mutation* operator is assigned

to make random changes in the solutions. A selection procedure, simulating the natural selection, selects a certain number of *parent* solutions, which the crossover uses to generate new solutions, also called *offsprings*. At the end of each iteration the offsprings together with the solutions from the previous generation form a new generation, after undergoing a selection process to keep a constant population size. The solutions are evaluated in terms of their *fitness values* identical to the fitness of individuals.

The GA are adaptive learning heuristic and they are generally referred to in plural, because several versions exist that are adjustments to different problems. They are also robust and effective algorithms that are computationally simple and easy to implement. The characteristics of GA that distinguishes them from the other heuristics, are the following [16]:

- GA work with coding of the solutions instead of the solution themselves. Therefore, a good, efficient representation of the solutions in the form of a chromosome is required.
- They search from a set of solutions, different from other metaheuristics like Simulated annealing and Tabu search that start with a single solution and move to another solution by some transition. Therefore they do a multi directional search in the solution space, reducing the probability of finishing in a local optimum.
- They only require objective function values, not e.g. continuous searching space or existence of derivatives. Real life examples generally have discontinuous search spaces.
- GA are nondeterministic, i.e. they are stochastic in decisions, which makes them more robust.
- They are blind because they do not know when they have found an optimal solution.

2.2.2 The Algorithm for VRP

As written above, GA easily adapts to different problems so there are many different versions depending on the problem to solve. There are, among other things, several ways to maintain a population and many different operators can be applied. But all GA must have the following basic items that need to be carefully considered for the algorithm to work as effective as possible [14]:

- A good genetic representation of a solution in a form of a chromosome.
- An initial population constructor.
- An evaluation function to determine the fitness value for each solution.
- Genetic operators, simulating reproduction and mutation.
- Values for parameters; population size, probability of using operators, etc.

A good representation or coding of VRP solution must identify the number of vehicles, which customers are assigned to each vehicle and in which order they are visited. Sometimes solutions are represented as binary strings, but that kind of representation does not suit VRP well. It is easy to specify the number of vehicles and which customers are inside each vehicle but it becomes too complicated when the order of the customers needs to be

given. Using the numeration of the customers instead, solves that problem. A suitable presentation of solutions to VRP is i.e. a chromosome consisting of several routes, each containing a subset of customers that should be visited in the same order as they appear. Every customer has to be a member of exactly one route. In figure 2.2 an example of the representation is shown for the solution in figure 2.1.

$$1: \boxed{5 \ 1 \ 3 \ 6} \qquad 2: \boxed{9 \ 7 \ 2} \qquad 3: \boxed{8 \ 10 \ 4}$$

Figure 2.2: A suitable representation of a potential VRP solution.

The construction of the initial population is of great importance to the performance of GA, since it contains most of the material the final best solution is made of. Generally, the initial solutions are randomly chosen, but they can also be results of some construction methods. It is called *seeding* when solutions of other methods join the randomly chosen solutions in the population.However, one should be careful to use too good solution at the beginning because those solutions can early become too predominant in the population. When the population becomes too homogeneous the GA loses its ability to search the solution space until the population slowly gains some variation by the mutation.

Recently, researchers have been making good progress with parallel GA, using multiple populations or subpopulations that evolve independently using different versions of GA. However, this project uses a sequential version with only one population. The *population size M* affects the performance of GA as well as affecting the convergence rate and the running time [16]. Too small population may cause poor performance, since is does not provide enough variety in the solutions. A large M usually provides better performance avoiding premature convergence. The convergence is discussed in section 2.2.4. The population size is definitely among the parameters that need tuning in order to find the value suitable for each problem. Although a constant population is used here, it is also possible to use a dynamic population, reducing the population size as the number of iterations increases. It has been experimented that the most rapid improvements in the population occur in the early iterations [16]. Then the changes become smaller and smaller and at the same time the weaker individuals become decreasingly significant.

In each iteration a number of parent solutions is selected and a crossover and/or other operators are applied producing offsprings. Maintaining the populations can be done in two ways. Firstly, by first selecting the new population from the previous one and then apply the operators. The new population can either include both "old" solutions from the previous population and offsprings or only offsprings, depending on the operators. Secondly, the operators can be applied first and then the new population is selected from both "old" solutions and offsprings. In order to keep a constant population size, clearly some solutions in the previous population will have to drop out. The algorithms can differ in how large proportion of the population is replaced in each iteration. Algorithms that replace a large proportion of the population are called *generational* and those replacing a single solution or only few are called *steady-state* [22]. In this project a steady-state algorithm is used. Below a pseudocode for a typical steady-state algorithm is shown.

Steady-state()

Population(M)

while the stopping criterion is not satisfied **do**

    P1, P2 ← ParentsSelection(Population)

    O1 ← Crossover(P1,P1)

    O2 ← Mutation(O1)

    R ← SolutionOutSelection(Population)

    Replace(O2,R)

end while

The function Population(M) generates M random solutions. The two selection methods need to be more specified. Many selection methods are available for choosing both individuals to reproduce and also for surviving at the end of every iteration. The same parents can be chosen several times to reproduce. The selection methods use fitness values associated with each solution to compare the solutions. A further discussion of the selection methods is given in section 2.2.4 below and the evaluation of the fitness value is discussed in next section. Since this is a steady-state algorithm, a crossover can be applied in every generation because a large part of the population will always be preserved in the next generation. Other operators can also be applied after or instead of Mutation. The Replace function replaces individual R in the population with the offspring O2 in order to keep the size of the population constant. Of course, it is not wise to replace the best individual in the population.

2.2.3   The Fitness Value

In order to perform a natural selection every individual i is evaluated in terms of its *fitness value* f_i, determined by an evaluation function. The fitness value measures the quality of the solutions and enables them to be compared. In section 2.2.4, different selection methods are discussed considering selective pressure. Selecting individuals for both reproduction and surviving has a crucial effect on the efficiency of GA. Too greedy a selection will lead to a premature convergence, which is a major problem in GA [14]. Since the selection methods are based on the fitness values, it is important to choose the evaluation function carefully.

Premature convergence can also be avoided by *scaling* the fitness values [16]. Scaling can be useful in later runs when the average fitness of the population has become close to

the fitness of the optimal solution and thus the average and the best individuals of the population are almost equally likely to be chosen. Naturally, the evaluation function and scaling of fitness values work together. Several scaling methods have been introduced, e.g. linear scaling, with and without sigma truncation and power law scaling [14].

The *linear scaling* method scales the fitness value f_i as follows:

$$f_i' = a \times f_i + b \tag{2.8}$$

where a and b are chosen so that the average initial fitness and the scaled fitness are equal. The linear scaling method is quite good but it runs into problems in later iterations when some individuals have very low fitness values close to each other, resulting in negative fitness values [14]. Also, the parameters a and b depend only on the population but not on the problem.

The *sigma truncation* method deals with this problem by mapping the fitness value into a modified fitness value f_i'' with the following formula:

$$f_i'' = f_i - (\overline{f} - K_{mult} \times \sigma) \tag{2.9}$$

K_{mult} is a multiplying constant, usually between 1 and 5 [14]. The method includes the average fitness $\overline{f}$ of the population and the standard deviation σ, which makes the scaling problem dependent. Possible negative values are set equal to zero. The linear scaling is now applied with f_i'' instead of f_i'.

Finally, there is the *power law scaling* method, which scales the fitness value by raising it to the power of k, depending on the problem.

$$f_i' = f_i^k \tag{2.10}$$

Often, it is straightforward to find an evaluation function to determine the fitness value. For many optimisation problems the evaluation function for a feasible solution is given, i.e. for both TSP and VRP, the most obvious fitness value is simply the total cost or distance travelled. However, this is not always the case, especially when dealing with multi objective problems and/or infeasible solutions.

There are two ways to handle infeasible solutions; either rejecting them or penalising them. Rejecting infeasible solutions simplifies the algorithm and might work out well if the feasible search space is convex [14]. On the other hand, it can have some significant limitations, because allowing the algorithm to cross the infeasible region can often enable it to reach the optimal solution.

Dealing with infeasible solutions can be done in two ways. Firstly, by extending the searching space over the infeasible region as well. The evaluation function for an infeasible solution $eval_u(x)$ is the sum of the fitness value of the feasible solution $eval_f(x)$ and either the penalty or the cost of repairing an infeasible individual $Q(x)$, i.e.

$$eval_u(x) = eval_f(x) \pm Q(x) \tag{2.11}$$

Designing the penalty function is far from trivial. It should be kept as low as possible without allowing the algorithm to converge towards infeasible solutions. It can be difficult to find the balance in between. Secondly, another evaluation function can be designed, independent of the evaluation function for the feasible solution $eval_f$.

Both methods require a relationship between the evaluation functions established, which is among the most difficult problems when using GA. The relationship can either be established using an equation or by constructing a global evaluation function:

$$eval(x) = \begin{cases} q_1 \cdot eval_f(x) & if \quad x \in \mathcal{F} \\ q_2 \cdot eval_u(x) & if \quad x \in \mathcal{U} \end{cases} \tag{2.12}$$

The weights q_1 and q_2 scale the relative importance of $eval_f$ and $eval_u$ and $\mathcal{F}$ and $\mathcal{U}$ denote the feasible region and the infeasible region respectively.

The problem with both methods is that they allow an infeasible solution to have a better fitness value than a feasible one. Thus, the algorithm can in the end converge towards an infeasible final solution. Comparing solutions can also be risky. Sometimes it is not quite clear whether a feasible individual is better than an infeasible one, if an infeasible individual is extremely close to the optimal solution. Furthermore, it can be difficult to compare two infeasible solutions. Consider two solutions to the *0-1 Knapsack problem*, where the objective is to maximise the number of items in the knapsack without violating the weight constraint of 99. One infeasible solution has a total weight of 100 consisting of 5 items of weight 20 and the other one has the total weight 105 divided on 5 items but with one weighing 6. In this specific situation the second solution is actually "closer" to attaining the weight constraint than the first one.

2.2.4   Selection

It seems that the population diversity and the selective pressure are the two most important factors in the genetic search [14]. They are strongly related, since an increase in the selective pressure decreases the population diversity and vice versa. If the population becomes too homogeneous the mutation will almost be the only factor causing variation in the population. Therefore, it is very important to make the right choice when determining a selection method for GA.

A selection mechanism is necessary when selecting individuals for both reproducing and surviving. A few methods are available and they all try to simulate the natural selection, where stronger individuals are more likely to reproduce than the weaker ones. Before discussing those methods, it is explained how the selective pressure influences the convergence of the algorithm,

Selective pressure

A common problem when applying GA, is a premature or rapid convergence. A convergence is a measurement of how fast the population improves. Too fast improvement

indicates that the weaker individuals are dropping out of the population too soon, i.e. before they are able to pass their characteristics on. The selective pressure is a measurement of how often the top individuals are selected compared to the weaker ones. Strong selective pressure means that most of the time top individuals will be selected and weaker individuals will seldom be chosen. On the other hand, when the selective pressure is weak, the weaker individuals will have a greater chance of being selected.

Figure 2.3 illustrates this for a population of five solutions with fitness values according to the size of its quadrangle. The y-axis shows the probability for each solution of being chosen. The line sp1 shows a strong selective pressure, where the top solutions are much more likely to be chosen than the weaker ones and line sp2 shows weaker selective pressure where the difference between the probabilities of selecting the solutions is smaller.

Strong selective pressure encourages rapid convergence but, on the other hand, too weak selective pressure makes the search ineffective. Therefore, it is critical to balance the selective pressure and the population diversity to get as good solution as possible.

Figure 2.3: Selective pressure.

Roulette Wheel Method

Firstly, there is a proportional selection process called, the *Roulette Wheel*, which is a frequently used method. In section 2.2.3, it is explained how every individual is assigned a fitness value indicating its quality. In the roulette wheel method, the probability of choosing an individual is directly proportional to its fitness value.

Figure 2.4 illustrates the method in a simple way for a problem having five individuals in a population. Individual P1 has a fitness value $f1$, P2 has $f2$, etc. Considering a pin at the top of the wheel, one can imagine when spinning the wheel that it would most frequently point to individual P3 and that it in the fewest occasions would point to individual P4. Consequently, the one with the largest fitness value becomes more likely to be selected as a parent than one with a small fitness value.

```
f1 = 0.28
f2 = 0.18
f3 = 0.33
f4 = 0.06
f5 = 0.15
```


Figure 2.4: Roulette Wheel method.

The drawback of the Roulette Wheel Method is that it uses the fitness values directly. That can cause some problems e.g. when a solution has a very small fitness value compared to the others, resulting in very low probability of being chosen. The ranking method in next chapter has a different approach.

Ranking

The second method is the *Ranking* method, which has been giving improving results [16]. It provides a sufficient selective pressure to all individuals by comparing relative goodness of the individuals instead of their actual fitness values. It has been argued that in order to obtain a good solution using GA, an adequate selective pressure has to be maintained on all the individuals by using a relative fitness measure [16]. Otherwise, if the population contains some very good individuals, they will early on become predominant in the population and cause a rapid convergence.

In Ranking, the individuals are sorted in ascending order according to their fitness. A function depending on the rank is used to select an individual. Thus it is actually selected proportionally to its rank instead of its fitness value as in the roulette wheel method. For instance, the selection could be based on the probability distribution below.

$$p(k) = \frac{2k}{\mathbf{M}(\mathbf{M}+1)} \tag{2.13}$$

The constant k denotes the kth individual in the rank and M is the size of the population. The best individual $(k = M)$ has a probability $\frac{2}{\mathbf{M}+1}$ of being selected and the worst individual $(k = 1)$ has $\frac{2}{\mathbf{M}(\mathbf{M}+1)}$ of being selected. The probabilities are proportional depending on the population size instead of fitness value.

The advantage of the Ranking method is that it is better able to control the selective pressure than the Roulette Wheel method. There are though also some drawbacks. The method disregards the relative evaluations of different solutions and all cases are treated uniformly, disregarding the magnitude of the problem.

Tournament Selection

The *Tournament Selection* is an efficient combination of selection and ranking methods. A parent is selected by choosing the best individual from a set of individuals or a subgroup from the population. The steady-state algorithm on page 21 requires only two individuals for each parent in every iteration and a third one to be replaced by the offspring at the end of the iteration. The method is explained considering the steady-state algorithm.

At first, two subgroups of each S individuals are randomly selected, since two parents are needed. If k individuals of the population were changed in each iteration, the number of subgroups would be k. Each subgroup must contain at least two individuals, to enable a comparison between them. The size of the subgroups influences the selective pressure, i.e. more individuals in the subgroups increase the selection pressure on the better individuals. Within each subgroup, the individuals compete for selection like in a tournament. When selecting individuals for reproduction the best individual within each subgroup is selected. On the other hand, the worst individual is chosen when the method is used to select a individual to leave the population. Then the worst individual will not be selected for reproduction and more importantly the best individual will never leave the population.

The Tournament Selection is the selection method that will be used in this project for both selection of individuals for reproduction and surviving. It combines the characteristics of the Roulette Wheel and the Ranking Method and is without the drawbacks of these methods have.

2.2.5   Crossover

The main genetic operator is *crossover*, which simulates a reproduction between two organisms, the parents. It works on a pair of solutions and recombines them in a certain way generating one or more offsprings. The offsprings share some of the characteristics of the parents and in that way the characteristic are passed on to the future generations. It is not able to produce new characteristics.

The functionality of the crossover depends on the data representation and the performance depends on how well it is adjusted to the problem. Many different crossover operators have been introduced in the literature. In order to help demonstrating how it works, the *Simple Crossover* [16] is illustrated in figure 4.1. The illustration is made with binary data presentation, even though it will not be used further in this project.

The Simple Crossover starts with two parent solutions P1 and P2 and chooses a random cut, which is used to divide both parents into two parts. The line between customers no. 2 and 3 demonstrates the cut. It generates two offsprings O1 and O2 that are obtained by putting together customers in P1 in front of the cut and customers in P2 after the cut and vice versa.

P1: `1 0 1 1 0 1`

P2: `1 1 1 0 1 1`

O1: `1 1 1 1 0 1`   O2: `1 0 1 0 1 1`

Figure 2.5: Illustration of Simple Crossover. The offspring O1 is generated from the right half of P1 and the left half of P2 and O2 is made from the left half of P1 and the right half of P2.

2.2.6   Mutation

Another operator is *mutation*, which is applied to a single solution with a certain probability. It makes small random changes in the solution. These random changes will gradually add some new characteristics to the population, which could not be supplied by the crossover. It is important not to alter the solutions too much or too often because then the algorithm will serve as a random search. A very simple version of the operator is shown in figure 2.6.

↓

P: `1 0 1 1 0 1`

O: `1 1 1 1 0 1`

Figure 2.6: Illustration of a simple mutation. A bit number 2 has been changed from 0 to 1 in the offspring.

The binary data string P represents a parent solution. Randomly, the second bit has been chosen to be mutated. The resulting offspring O illustrates how the selected bit has been changed from 0 to 1.

2.2.7   Inversion

The third operator is *Inversion*, which reverses the order of some customers in a solution. Similar to the mutation operator, it is applied to a single solution at a time. In figure 2.7 this procedure is illustrated with a string of letters, which could represent a single route in solution.

Two cuts are randomly selected between customers 3 and 4 and 7 and 8, respectively. The order of the customers between the cuts is reversed.

The inversion operator will not be used specifically in this project. However, the Local Search Algorithms in the next chapter reverse the order of the customers in a route if it improves the solution.

| h d c | j f e a | g b i | $\longrightarrow$ | h d c | a e f j | g b i |

Figure 2.7: A single route before(left) and after(right) an inversion. The order of the letters between the lines has been reversed.

2.3   Summary

In this chapter the Vehicle Routing Problem has been described. The basic concepts of Genetic Algorithms were introduced, such as the fitness value, the crossover and the mutation operators. In the next chapter the development of the Local Search Algorithms will be explained.

Chapter 3

Local Search Algorithms

The experience of the last few years has shown that combining Genetic Algorithms with *Local Search Algorithms* (LSA) is necessary to be able to solve VRP effectively [10]. The LSA can be used to improve VRP solutions in two ways. They can either be improvement heuristics for TSP that are applied to only one route at a time or multi-route improvement methods that exploit the route structure of a whole solution [13]. In this project, LSA will only be used to improve a single route at a time.

Most local search heuristics for TSP can be described in a similar way as Lin's λ-*Opt* algorithm [12]. The algorithm removes λ edges from the tour and the remaining segments are reconnected in every other possible way. If a profitable reconnection is found, i.e. the first or the best, it is implemented. The process is repeated until no further improvements can be made and thus a *locally optimal* tour has been obtained. The most famous LSA are the simple *2-Opt* and *3-Opt* algorithms ($\lambda=2$ and $\lambda=3$). The 2-Opt algorithm, which was first introduced by Croes in 1958 [1], removes two edges from a tour and reconnects the resulting two subtours in the other possible way. Figure 3.1 is an illustration of a single step in the 2-Opt algorithm. The illustration is only schematic (i.e. if the lengths were as they are shown, this move would not have been implemented). For simplicity later on, the tour is considered directed.

Figure 3.1: A tour before (left) and after (right) a 2-Opt move.

The 3-Opt algorithm was first proposed by Bock in 1958 [1]. It deletes three edges from a tour and reconnects the three remaining paths in some other possible way. The 3-Opt

algorithm is not implemented here because it is not likely to pay off. This is shown in [1] where test results propose that for problems of 100 customers the performance of 3-Opt is only 2% better than 2-Opt. The biggest VRP that will be solved in this project has 262 customers and minimum 25 vehicles (see chapter 7) thus each route will most likely have considerably fewer customers than 100. Therefore, the difference in performance can be assumed to be even less. Furthermore, 3-Opt is more time consuming and difficult to implement.

There are different ways to make both 2-Opt and 3-Opt run faster. For instance by implementing a *neighbour-list*, which stores the k nearest neighbours for each customer [1]. As an example, consider a chosen t1 and t2. The number of possible candidates for t3 (see figure 3.1) is reduced to k instead of $n - 3$ where n is the number of customers in the route. However, since the algorithm will be applied to rather short routes, as was explained above, it will most likely not pay off. The emphasis will be on producing rather simple but effective and 2-Opt algorithms. The 2-Opt algorithm is very sensitive to the sequence in which moves are performed [11]. Considering the sequence of moves three different 2-Opt algorithms have been put forward. In the following sections they are explained and compared. The best one will be used along in the process.

3.1   Simple Random Algorithm

The *Simple Random Algorithm* (SRA) is the most simple 2-Opt algorithm explained in this chapter. It starts by randomly selecting a customer t1 from a given tour, which is the starting point of the first edge to be removed. Then it searches through all possible customers for the second edge to be removed giving the largest possible improvement. It is not possible to remove two edges that are next to each other, because that will only result in exactly the same tour again. If an improvement is found, the sequence of the customers in the tour is rearranged according to figure 3.1. The process is repeated until no further improvement is possible.

```
Simple Random(tour)

savings ← 1

while savings > 0 do
        t1ind ← random(0, length[tour]-1)
        t1 ← tour[t1ind]
        t2ind ← t1ind+1 mod length[tour]
        t2 ← tour[t2ind]
        savings ← 0

        for tf ← 0 to length[tour]-1
            if tf ≠ t1ind AND tf ≠ t2ind AND tf+1 mod length[tour] ≠ t1ind
                t4ind ← tf
                t4 ← tour[t4ind]
                t3ind ← t4ind + 1 mod length[tour]
                t3 ← tour[t3ind]
                distanceDiff ← dist[t1][t2]+dist[t4][t3]-dist[t2][t3]-dist[t1][t4]
                if distanceDiff > savings
                    savings ← distanceDiff
                    fint3 ← t3ind
                    fint4 ← t4ind
        end for

        if savings > 0
            Rearrange(t1ind, t2ind, fint3, fint4)

end while
```

An obvious drawback of the algorithm is the choice of t1, because it is possible to choose the same customer as t1, repeatedly. The algorithm terminates when no improvement can be made using that particular t1, which was selected at the start of the iteration. However, there is a possibility that some further improvements can be made using other customers as t1. Thus, the effectivity of the algorithm depends too much on the selection of t1. The algorithm proposed in next section handles this problem by not allowing already selected customers to be selected again until in next iteration.

3.2   Non Repeating Algorithm

The *Non Repeating Algorithm* (NRA) is a bit more complicated version of the Simple Random algorithm. A predefined selection mechanism is used to control the random selection of t1, instead of choosing it entirely by random. The pseudocode for the algorithm

is shown below.

```
Non Repeating(tour)

savings ← 1

while savings > 0 do
      selectionTour ← tour
      top ← length[selectionTour]-1
      savings ← 0

      for t ← 0 to length[selectionTour]-1
          selind ← random(0, top)
          (t1, t1ind) ← findInTour(selectionTour[selind])
          exchange selectionTour[top] ↔ selectionTour[selind]
          t2ind ← t1ind+1 mod length[tour]
          t2 ← tour[t2ind]
          savings ← 0

          for tf ← 0 to length[tour]
              if tf ≠ t1ind AND tf ≠ t2ind AND tf+1 mod length[tour] ≠ t1ind
                  t4ind ← tf
                  t4 ← tour[t4ind]
                  t3ind ← t4ind + 1 mod length[tour]
                  t3 ← tour[t3ind]
                  distanceDiff ← dist[t1][t2]+dist[t4][t3]-dist[t2][t3]-dist[t1][t4]
                  if distanceDiff > savings
                     savings ← distanceDiff
                     fint3 ← t3ind
                     fint4 ← t4ind
          end for

          if savings > 0
              Rearrange(t1ind, t2ind, fint3, fint4)
      end for
end while
```

The selection mechanism is implemented in the outmost for loop. It allows each customer in the tour to be selected only once in each iteration (inside the while-loop). The customers are randomly selected one by one and when they have been used as t1, they are eliminated from the selection until in next iteration. Figure 3.2 shows a single step using the technique.

```
                          sel

                           │
                           ▼
                    ┌───────────┐            ┌───────────┐
                    │ 3 5 2 1 4 │   ────▶    │ 3 4 2 1 5 │
                    └───────────┘            └───────────┘
                           ▲                      ▲
                           │                      │

                          top                    top
```

Figure 3.2: Selection mechanism for t1. The first customer *sel* is selected randomly among the five customers. Afterwards, it switches places with the last customer and the pointer *top* is reduced by one. The second customers is selected among the four customers left.

Considering the tour at the left hand side in the figure the process is following: Firstly, a pointer *top* is set at the last customer. Secondly, customer no. 5 is randomly chosen from the customers having indices 1 to *top*. Then customer no. 5 and the one being pointed at, which is customer no. 4, switch places. Finally, the pointer is reduced by one, so in next step customer no. 5 has no possibility of being chosen again in this iteration.

In the beginning of each iteration the algorithm starts by making a copy of the tour into *selectionTour*, in order to preserve the original tour. Then t1 is randomly selected and edge e1 is defined. By going through all the potential customers in the tour, the customer t4 providing the best improvement is found. As in SRA, it is disallowed to choose the second edge e2 next to e1 because that will generate the same tour again. If an improvement to the tour is found, the best one is implemented. Otherwise the algorithm terminates. The final iteration does not improve the tour but it is necessary to verify that no further improvements can be made. When termination occurs a local optimum tour has been found.

3.3   Steepest Improvement Algorithm

The *Steepest Improvement Algorithm* (SIA) has a bit different structure than the two previous 2-opt algorithms. SRA and NRA choose a single customer t1, find the customer t4 among other customers in the tour that will give the largest saving and rearrange the tour. SIA, on the other hand, compares all possible combinations of t1 and t4 to find the best one and then the tour is rearranged. This means that it performs more distance evaluations for each route rearrangement. Each time the largest saving for the tour is performed. The algorithm is best explained by the following pseudocode.

Steepest Improvement(tour)

savings $\leftarrow$ 1

while savings > 0 **do**
    savings $\leftarrow$ 0
    for t1ind $\leftarrow$ 0 **to** length[tour]-1
        for t4ind $\leftarrow$ 0 **to** length[tour]-1
            if t4ind $\neq$ t1ind AND t4ind $\neq$ t1ind+1 AND t4ind+1$\neq$ t1ind
                t1 $\leftarrow$ tour[t1ind]
                t2 $\leftarrow$ tour[t1ind+1]
                t3 $\leftarrow$ tour[t4ind+1]
                t4 $\leftarrow$ tour[t4ind]
                distanceDiff $\leftarrow$ distance[t1][t2]+distance[t4][t3]-distance[t2][t3]
                -distance[t1][t4]
                if distanceDiff>savings
                    savings $\leftarrow$ distanceDiff
                    t1best $\leftarrow$ t1ind
        end for
    end for
    if savings > 0
        Rearrange(t1best,t1best+1,t4best+1,t4best)

end while

There is no randomness involved in the selection of t1. Every combination of t1 and t4 is tested for possible improvements and the one giving the largest improvement is implemented. It is necessary to go through all possibilities in the final iteration to make sure that no further improvements can be made.

3.4   The Running Time

It is very difficult to estimate the running time of the algorithms theoretically. As was written on page 30, the algorithms are very sensitive to the sequence in which the moves are performed. Naturally, the running time depends on the problem but it also depends on the original solution. It is particularly hard to estimate the running time of SRA and NRA, where the selection sequence is based on random decisions.

However, the relative running time of the operators can be estimated by the means of their structure. In both SRA and NRA, a rearrangement of a tour is made after almost n comparisons. On the other hand, each rearrangement of a tour in SIA requires a little less than n^2 comparisons. It is therefore expected that SRA and NRA have similar running times and compared to them, SIA has longer running time.

3.5   Comparison

Before carrying on, one of the Local Search Algorithms is chosen for further use in the project. The performance of the algorithms is only compared for relatively small problems with 50 customers at most. The largest problem, which GA is applied to in this project has 262 customers (see chapter 7) therefore it is fair to assume that the routes will not have more customers than 50. Ten problems were generated using the 5, 10, 15, 20, 25, 30, 35, 40, 45 and 50 first customers in problem kroD100 from [19]. The problems were solved to optimality by Thomas Stidsen [18] using a branch and cut algorithm. The values are shown in appendix A. The algorithms were run 5 times on each of the ten instances and the difference from optimum, standard deviation and time was recorded. Table 3.1 shows the results.

| Problem | SRA | | | NRA | | | SIA | | |
| | Diff. | Std. dev. | Time | Diff. | Std. dev. | Time | Diff | Std. dev. | Time |
sizes	(%)	σ	(ms)	(%)	σ	(ms)	(%)	σ	(ms)
5	1,67	1,45	36	0,00	0,00	20	0,00	0,00	20
10	18,54	14,66	18	0,48	0,49	16	0,48	0,49	14
15	58,80	30,45	16	5,33	4,36	18	6,76	6,00	22
20	77,87	46,63	28	2,99	2,96	32	5,52	1,57	26
25	97,87	75,47	10	9,50	2,31	12	8,15	4,13	12
30	109,08	30,54	14	6,64	4,79	14	4,31	4,14	18
35	138,14	36,95	10	6,25	4,01	10	4,69	2,87	20
40	143,32	79,61	18	7,20	2,75	18	7,45	4,36	42
45	121,23	37,71	16	9,24	5,12	16	5,40	3,51	36
50	118,10	24,37	14	5,08	1,32	16	5,85	2,59	46
Average	88,40	37,78	18	5,27	2,81	17	4,86	2,97	26

Table 3.1: The performance of the Local Search Algorithms. SRA is outperformed by NRA and SIA. NRA and SIA both perform quite well but the average difference from optimum is smaller for SIA.

The percentage difference from optimum is plotted in a graph in figure 3.3. Figure 3.4 illustrates how the cost gradually improves with the number of iterations. The data is collected during a single run of each of the algorithms when solving the problem with 25 customers.

Figure 3.3: Percentage difference for SRA, NRA and SIA. SRA is clearly outperformed by NRA and SIA, which perform almost equally well. SIA gives a bit better results.

Figure 3.4: The development of the cost for SRA, NRA and SIA. SRA is clearly not effective enough. SIA converges slower towards the optimal value than NRA but it gets a little bit closer to it.

It is quite clear that SRA is not able to compete with neither NRA nor SIA. The difference

from optimum is much larger, even though the time it uses is relatively short. The difference from optimum is a little bit smaller for SIA compared to NRA, but the time is considerably worse. In the latter figure it is illustrated how the convergence of SIA is much slower than of SRA and NRA and it requires more iterations to obtain a good solution.

SIA is chosen to be used in the project. According to the results, it provides a little bit better results and that is considered more important than the time. When the Local Search Algorithm of choice is applied with other genetic operators in the final testing it is believed that they account for most of the time therefore the choice is mainly based on the difference from optimum.

3.6 Summary

In this chapter, three Local Search Algorithm were developed; Simple Random Algorithm, Non Repeating Algorithm and Steepest Improvement Algorithm. Steepest Improvement Algorithm was chosen to use further in the project, based on its accuracy. The next chapter describes the main part of the project, which involves the development of the fitness value and the genetic operators.

Chapter 4

The Fitness Value and the Operators

The genetic operators and the evaluation function are among the basic items in GA (see page 19). The operators can easily be adjusted to different problems and they need to be carefully designed in order to obtain an effective algorithm.

The geography of VRP plays an essential role when finding a good solution. By the geography of a VRP it is referred to the relative position of the customers and the depot. Most of the operators that are explained in this chapter take this into consideration. The exceptions are Simple Random Crossover and Simple Random Mutation, which depend exclusively on random choices. They were both adopted from the original project, see chapter 1. Some of the operators are able to generate infeasible solutions, with routes violating the capacity constraint, thus the fitness value is designed to handle infeasible solutions.

Before the fitness value and the different operators are discussed, an overview of the main issues of the development process is given.

Overview of the Development Process

1. The process began with designing three Local Search Algorithms that have already been explained and tested in chapter 3.

2. In the beginning, infeasible solutions were not allowed, even though the operators were capable of producing such solutions. Instead, the operators were applied repeatedly until they produced a feasible solution and first then the offspring was changed. That turned out to be a rather ineffective way to handle infeasible solutions. Instead the solution space was relaxed and a new fitness value was designed with an additional penalty term depending on how much the vehicle capacity was violated. This is explained in the next section.

3. The Biggest Overlap Crossover (see section 4.2.2) was the first crossover operator to be designed, since Simple Random Crossover was adopted from the previous project, see chapter 1. Experiments showed that both crossover operators were producing

offsprings that were far from being feasible, i.e. the total demand of the routes was far from being within the capacity limits. The Repairing Operator was generated to carefully make the solutions less infeasible, see section 4.4.1.

4. The Horizontal Line Crossover (see section 4.2.3) gave a new approach that was supposed to generate offsprings, which got their characteristics more equally from both parents. However, the offsprings turned out to have rather short routes and too many of them did not have enough similarity to their parents. Geographical Merge was therefore designed to improve the offsprings by merging short routes. The Horizontal Line Crossover is discussed in section 4.2.3 and Geographical Merge is considered in section 4.4.2.

5. Finally, Uniform Crossover was implemented. It was a further development of Horizontal Line Crossover, in order to try to increase the number of routes that were transferred directly from the parent solutions. The operator is explained in section 4.2.4.

4.1   The Fitness Value

Every solution has a fitness value assigned to it, which measures its quality. The theory behind the fitness value is explained in section 2.2.3. In the beginning of the project, no infeasible solutions were allowed, i.e. solutions violating the capacity constraint, even though the operators were able to generate such solutions. To avoid infeasible solutions the operators were applied repeatedly until a feasible solution was obtained, which is inefficient and extremely time consuming. Thus, at first the fitness value was only able to evaluate feasible VRP solutions.

It is rather straight forward to select a suitable fitness value for a VRP where the quality of a solution s is based on the total cost of travelling for all vehicles;

$$f_s = \sum_r cost_{s,r} \tag{4.1}$$

where $cost_{s,r}$ denotes the cost of route r in solution s.

Although it is the intention of GA to generate feasible solutions, it can often be profitable to allow infeasible solutions during the process. Expanding the search space over the infeasible region does often enable the search for the optimal solution, particularly when dealing with non-convex feasible search spaces [16], as the search space of large VRP. The fitness value was made capable of handling infeasible solutions by adding a penalty term depending on how much the capacity constraint is violated. The penalty was supposed to be insignificant at the early iterations, allowing infeasible solutions, and predominant in the end to force the the final solution to be feasible. Experiments were needed to find the right fitness value that could balance the search between infeasible and feasible solutions.

It is reasonable to let the penalty function depend on the number of iterations, since it is supposed to develop with increasing number of iterations. The exponential function depending on the number of iteration $\exp(it)$ was tried, since it had just the right form. Unfortunately, in the early iterations the program ran into problems because of the size of the penalty term. The program is implemented in Java and the biggest number Java can handle is approx. 92234×10^{18}. Already in iteration 44, the penalty function grew beyond those those limits $(\ln(92234 \times 10^{18}) = 43.6683)$. It also had the drawback that is did not depend on the problem at all and it always grew equally fast no matter how many iterations were supposed to be performed.

A new more sophisticated evaluation function for the fitness value was then developed. It is illustrated in equations 4.2 to 4.4.

$$f_s = \sum_{r \in s} cost_{s,r} + \alpha \cdot \frac{it}{IT} \sum_{r \in s} (max(0, totdem_{s,r} - cap))^2 \tag{4.2}$$

$$\alpha = \frac{best}{\frac{1}{IT}(\frac{mnv}{2} \cdot cap)^2} \tag{4.3}$$

$$mnv = \left\lceil \frac{\sum_{c \in s} dem_c}{cap} \right\rceil \tag{4.4}$$

where:

it	is the current iteration,
IT	denotes the total number of iterations,
$totdem_{r,s}$	is the total demand of route r in solution s,
cap	represents the uniform capacity of the vehicles,
$best$	is the total cost of the best solution in the beginning and
dem_c	denotes the demand of customer $c \in s$.

The left part of the evaluation function is just the cost as in equation 4.1. It denotes the fitness value of a feasible solution because the second part equals zero if the capacity constraint is attained. The second part is the penalty term. The quantity of the violation of the capacity constraint is raised to the power of 2 and multiplied with a factor α and the relative number of iterations. By multiplying with $\frac{it}{IT}$ the penalty factor is dependent on where in the process it is calculated, instead of the actual number of the current iteration. The factor α makes the penalty term problem dependent, because it includes the cost of the best solution in the initial population. It also converts the penalty term into the same units as the first part of the evaluation function has.

The size of the α determines the effect of the penalty, i.e. a large α increases the influence of the penalty term on the performance and a small α decreases the effect. Figures 4.1 to 4.4 show four graphs that illustrate the development of the total cost of the best individual and the total cost, the cost and the penalty of the offspring for three different values of

α. Since α is calculated by the means of equation 4.3, the graphs show the effect of multiplying α with a scalar. The results were obtained by solving the problem instance A-n80-k10 from [20] with Simple Random Crossover, Simple Random Mutation (pSame = 30%, rate = 100%), Repairing Operator (rate = 100 %) and Steepest Improvement Algorithm, which are all explained in the following sections. The population size was 50 and the number of iterations was 10000.

Figure 4.1: The development of the cost of the best individual for three different values of α as the number of iterations increases.

Figure 4.2: The development of the total cost of the offspring for three different values of α as the number of iterations increases.

Figure 4.3: The development of the cost of the offspring for three different values of α as the number of iterations increases.

Figure 4.4: The development of the penalty of the offspring for three different values of α as the number of iterations increases.

The main purpose of the graphs in the figures above is to show how the size of the penalty varies for different values of α and the relative size to the cost. In figures 4.1 and 4.3 it seems as if the convergence becomes more rapid as the value of α decreases, although the difference is really small. Even though there is a difference is the convergence the final results are almost identical. As always it is risky to jump to conclusions based on a single problem because the convergence also depends on the shape of the solutions space. Different population sizes or operators will probably affect the convergence more. Figures 4.2 and 4.4 illustrate how both the penalty and the total cost of the offspring gradually increases with the number of iterations and as the value of α increase the penalty also increases significantly. The values of the y-axis show the size of the penalty compared to the cost. The graphs also show that most of the time the offspring represents an infeasible solution, but once in a while a feasible solution is generated since the total cost of the best individual gradually reduces.

4.2   The Crossover Operators

In chapter 2.2.5 the general idea behind the crossover operator was discussed and a very simple crossover was illustrated. In this chapter, four more complex crossover operators are introduced and their development phase is discussed when it is appropriate. All the crossover operators need two parent solutions P1 and P2 to generate one offspring. P1 is always the better solution.

4.2.1   Simple Random Crossover

The *Simple Random Crossover* (SRC) is the most simple one of the four crossover operators and it mostly depends on random decisions. In words, the operator randomly selects a subroute from P2 and inserts it into P1. The pseudocode is the following:

SRCrossover(P1, P2)

copy individual P1 into *offspring*

randomly select a *subroute* from P2

delete the members of the *subroute* from the *offspring*

bestInd $\leftarrow$ BestInsertion(*offspring*, *subroute*)

insertSubRoute(*offspring*, *bestInd*, *subroute*)

return *offspring*

At first, P1 is copied into the offspring, since both P1 and P2 are to be preserved in the population, until a decision is made to replace them. The offspring is modified and P1 remains untouched. Firstly, a route in P2 is randomly chosen and a subroute is selected from that particular route, also by random. The subroute contains at least one customer and at most the whole route. Before inserting the subroute into P1, all its customers are deleted from P1 to avoid duplications in the solution. It is more preferable to perform the deletion before the insertion, so the subroute can be inserted as a whole and left untouched in the offspring. At last, the subroute is inserted in the best possible place, which is found by the function BestInsertion. The function finds both the route in which the subroute is inserted and the two customers it is inserted between. Consider k_1 denoting the first customer in the subroute and k_n the last one and c_m and c_{m+1} being customers in a route in the offspring. The pay off of inserting the subroute between c_m and c_{m+1} is measured by the formula:

$$p_m = cost(c_m, c_{m+1}) - cost(c_m, k_1) - cost(k_n, c_{m+1}) \tag{4.5}$$

where $cost(c_m, c_{m+1})$ is the cost of travelling from c_m to c_{m+1}. The algorithm searches through the whole offspring and inserts the subroute in the place giving the largest payoff. A new offspring has been generated!

The operator can be described as unrefined. It does not consider the solutions it is working with at all, because all decisions are based on randomness. The subroute is inserted into P1, totally disregarding the capacity constraint of the vehicle. The insertion method can have some drawbacks, since it only looks at the first and the last customer in the subroute, which do not necessarily represent subroute as a whole. Also if P1 and P2 are

good solutions, they probably have almost or totally full vehicles on most of the routes and consequently the operator generates an infeasible solution. Since infeasible solutions are penalised, it can make the algorithms ineffective if SRC generates infeasible solutions most of the time.

Furthermore, the geography of the problem is ignored. If the subroute is chosen from a good or partially good route the operator does not make any effort to choose the good part to pass on to the offspring. A totally random selection of a subroute can overlook it or just take a part of it. As a consequence, too much use of random selection can make it difficult for good characteristics to spread out in the population. On the other hand, some randomness can be necessary to increase the diversity of the population.

In next chapter a crossover is introduced that compares the geography of the subroute to the geography of the offspring when inserting the subroute into the offspring.

4.2.2   Biggest Overlap Crossover

The *Biggest Overlap Crossover* (BOC) can be looked at as an extended version of SRC. It uses the geography of the solution, i.e. the relative position of the routes, in addition to the total demand of its routes, when inserting the subroute. Calculating the actual distance between every two routes can be complicated due to their different shapes. Therefore, so-called bounding boxes are used to measure the size of each route and to calculate the distance between them. Further explanation of bounding boxes is given below.

As in SRC, the subroute is randomly selected from P1. There are two possible approaches of taking the geography or capacity into consideration. The first one, starts by choosing three routes from P1 considering the size of the overlapping between the bounding boxes of the subroute and the routes of P1. The subroute is inserted into one of the three routes having the smallest total demand. The second approach first selects the three routes having the smallest total demand of the routes in P1, then the subroute is inserted into the one of the three routes having the biggest overlap with the subroute. Both approaches can generate infeasible solutions, if the subroute contains customers with too large total demand. The two approaches that are called *First Geography, then Capacity* (GC) and *First Capacity, then Geography* (CG) are discussed further below and a comparison is given.

Bounding boxes

Each route has its own *bounding box*, which is the smallest quadrangle the entire route fits in (the depot is also a member of every route). Figure 4.5 illustrates the bounding boxes for a solution with four routes.

Figure 4.5: Bounding boxes.

In order to estimate the distance between the routes the shortest distance between the bounding boxes of the routes is found. Often the bounding boxes will overlap, especially since all routes share the depot. In the figure, the two routes above the depot have overlapping bounding boxes. The size of the overlapping measures the closeness of the routes. Naturally, routes with overlapping bounding boxes are considered closer to each other than routes with non overlapping bounding boxes. If no bounding boxes overlap the routes with shortest distance between them are considered closest.

First Geography, then Capacity

At first the *First Geography, then Capacity* approach is discussed. The pseudocode is as follows:

BOCrossover(P1, P2)

copy individual P1 into *offspring*

randomly select a *subroute* from P2

determine the *bounding box* of the *subroute*

delete the members of the *subroute* from the *offspring*

biggestOverlap ← the 3 routes in P1 having the biggest overlap with the *subroute*

minDemInd ← the route in *biggestOverlap* with the smallest total demand

bestInd ← [*minDemInd*, BestInsertion(*offspring[minDemInd]*, *subroute*)]

insertSubRoute(*offspring*, *bestInd*, *subroute*)

return *offspring*

At first, individual P1 is copied into the offspring to preserve P1 in the population. Secondly, a subroute is randomly selected from P2 and its bounding box is calculated. Then its members are deleted from P1 to prevent duplications. By comparing the bounding boxes of the subroute to the bounding box of each route in P1 the three closest routes are determined. Finally, the subroute is inserted into the route with the smallest total demand, in the cheapest possible way, which is given by the function BestInsertion. The functionality of BestInsertion is described in the previous section.

First Capacity, then Geography

The second version *First Capacity, then Geography* (CG) considers first the capacity and then the distance between the routes. The process is very similar to the previous approach, except when it comes to choosing the insertion route in P1. Firstly, the three routes having the smallest total demands are chosen from individual P1. The subroute is inserted into one of the routes closest to the subroute. Again, the function BestInsertion determines the best place to insert the subroute into the route.

Comparison of GC and CG

The two methods GC and CG weigh the geography and the capacity differently. A comparison is made to find out if any considerable difference is between the performance of the two methods. That can give some indication of whether the emphasis should be on the geography or the capacity.

Table 4.1 shows the performance of the two approaches. The percentage difference from optimum, the standard deviation σ and time are stated for four problems of different sizes; A-n32-k5, A-n44-k6, A-n60-k9 and A-n80-k10 from [20]. The results show the average of 5 runs using Simple Random Crossover, Simple Random Mutation (pSame = 30% and rate = 100%) and Simple Improvement algorithm. The population size was 50 and 10000 iterations were performed.

| Problem | GC | | | CG | | |
sizes	Diff. from optimum (%)	Std. dev. σ	Time (ms)	Diff. from optimum (%)	Std. dev. σ	Time (ms)
32	5,08	2,45	3241	17,70	6,00	5125
44	22,08	4,04	5147	47,51	4,53	7168
60	23,40	4,77	7649	73,80	17,52	10876
80	33,56	5,12	12818	67,01	4,55	18490
Average	21,03	4,10	7214	51,51	8,15	10415

Table 4.1: Results from comparison of First Geography, then Capacity and First Capacity, then Geography.

In figure 4.6 the results are illustrated in a graph.

Figure 4.6: Comparison of GC (geocap) and CG (capgeo). CG is definitely outperformed by GC.

The results are clear. For all four problem instances the deviation from optimum is at least twice as large as for CG compared to GC. Besides, CG is more time consuming. Therefore, Biggest Overlap Crossover uses the GC approach in the rest of the project.

4.2.3   Horizontal Line Crossover

The *Horizontal Line Crossover* (HLC) introduces another method to generate offsprings. SRC and BOC generate offsprings that mostly come from P1, though it depends on the relative problem size to the vehicle capacity. The HLC tries to combine the solutions P1 and P2 more equally in the offspring. A horizontal line is drawn through the depot and divides both solutions into two parts. Bounding boxes (see section 4.2.2) are used to identify the routes. Routes in P1 having bounding boxes above the line are inserted untouched into the offspring along with routes in P2 that have bounding boxes below the line. This is illustrated in figure 4.7. For simplification, the figure shows the routes without the bounding boxes but in figure 4.5 solution P2 is illustrated with bounding boxes.

Figure 4.7: An illustration of the Horizontal Line Crossover. Routes in P1 above the line and route in P2 below the line are inserted untouched into the offspring. Note that the rightmost customer is not a member.

The line has to be drawn through the depot, since all routes contain the depot. It could also be possible to draw the line vertically. However, the algorithm would not have been near as effective if the line was drawn at an angle, since the bounding boxes are used to identify the routes. That would reduce the number of routes that come directly from the parent solutions into the offspring. In the ideal case, no routes cross the line and all customers are transfered directly from the parent solutions to the offspring. This is however not always the case and in order to include all the customers in the offspring another method is needed to gather the rest of the customers into the offspring. In figure 4.7 the rightmost customer is not yet a member of any route in the offspring.

All classical heuristics that belong to the construction heuristics or the two-phase methods can be used to generate new routes from the rest of the customers, see section 2.1.4. The two-phase method the Sweep Algorithm is chosen here. The idea behind the Sweep Algorithm first came forward in the early 70t's but generally it is attributed to Gillet and Miller that proposed it in 1974 [12]. In the first phase feasible routes are made by rotating a ray centred at the depot and gradually it gathers the customers into clusters or routes. In the second phase a TSP is solved for each route. Here the Steepest Improvement Algorithm from chapter 3 is used to improve the routes where the initial order of the customers is the order in which they are gathered. It can therefore be expected that HLC depends on Steepest Improvement Algorithm to provide good solutions.

The pseudocode is:

HLCrossover(P1, P2)

NRP1 ← number of routes in P1
NRP2 ← number of routes in P2
offjind ← 0

for n ← 0 **to** NRP1
    if P1[n] is above a horizontal line through the depot
        put route P1[n] into *offspring*[offjind]
        offjind ← offjind + 1
end for

for m ← 0 **to** NRP2
    if P2[m] is below a horizontal line through the depot
        put route P2[m] into *offspring*[offjind]
        offjind ← offjind +1
end for

put all other customers into *customersLeft*

sweep the customers in *customersLeft* into feasible routes and put into the *offspring*

return *offspring*

In order to generate a good crossover, it is essential that the generated offsprings reflect on its parents. Therefore, it must be the aim of the crossover to pass as many routes as possible from the parent solutions to the offspring, thus leaving as few customers as possible to the Sweep Algorithm. The disadvantage of HLC is that its quality depends on the relative position of the customers to the depot. If the depot is placed relatively high or low in the graph, the routes will mainly be obtained from only one parent solution. Also, if many customers are placed around the line, it is expected that many routes cross the line. Then too many customers will be put into the routes by the Sweep Algorithm and the crossover will basically become the Sweep Algorithm.

4.2.4   Uniform Crossover

The last type of crossover operator introduced is the *Uniform Crossover* (UC), which is named after the Uniform Crossover known in the literature, e.g. in [16]. Consider solutions represented as binary strings, then for each bit the crossover chooses between the bits in the parent solutions. In a similar way, UC chooses between the routes of the parent solutions based on the quality of the routes and as long as they do not conflict with each other. The goal is to design a crossover that passes more routes directly from the parents, compared to the HLC in the previous section, and which performance does

not depend so much on the position of the depots.

In order to be able to explain the algorithm, a small problem of 13 customers is illustrated in figure 4.8. It shows two parent solutions of the problem and the cost of their routes. All routes in both these solutions attain the capacity constraint, so there is no penalty involved in this case.


```
P1                                    P2

        1                             L1 = 16              LA = 20
             2                        L2 = 18        A        B    LB = 21
                                      L3 = 14                      LC = 14
             3                        L4 = 16        D    C        LD = 16
        4
```

Figure 4.8: Two parent solutions and their route costs.

When solving VRP, the objective is to minimise the total cost of visiting all the customers without violating the vehicle capacity constraint, see chapter 2.1. Thus, it is desirable to have the average cost of travelling between the customers in each route as low as possible. Therefore, the quotient $R_{s,r}$ is calculated for each route r of every solution s and used to measure the quality of them.

$$R_{s,r} = totCost_{s,r}/noCust_{s,r} \qquad (4.6)$$

The term $totCost_{s,r}$ is the sum of the cost and penalty for route r in solution s and $noCust_{s,r}$ denotes the number of customers in route r in solution s. The quotient for each route in the two solutions in figure 4.8 is:

P1	P2
$R_1 = 8{,}0$	$R_A = 5.0$
$R_2 = 6{,}0$	$R_B = 7.0$
$R_3 = 4{,}7$	$R_C = 4.7$
$R_4 = 3{,}2$	$R_D = 5.3$

Table 4.2: The quotient $R_{s,r}$ for each route in figure 4.8.

The idea is now to transfer as many untouched routes from the parents to the offspring. However, it is not possible to move two routes from different individuals if they conflict with each other, i.e. they share some customers. Each route in a solution will at least conflict with one route in the other individual. In order to choose between those routes, the quotient $R_{s,r}$ is used. The procedure is as follows; At first the route with the smallest

value of $R_{s,r}$ is chosen from P1 and all routes in P2 that conflict with it, are excluded. Then a route from P2 is selected among the non excluded routes and routes in P1 conflicting with it are eliminated. This process continues until either all routes have been selected or all the routes left have been eliminated. Ideally, all customers will be assigned to the offspring via the routes of the parent solutions. However, generally there will be some customers that are left out.

Figure 4.9 shows how the routes in the solutions in figure 4.8 conflict with each other. Firstly, route 4 is selected from P1 having the smallest quotient according to table 4.2, which eliminates routes A, C and D. Secondly, route B is selected, excluding routes 2 and 3. Thirdly, route 1 is chosen and the process terminates. The resulting offspring has three routes with altogether 10 customers. The last three customers need to be inserted into the offspring by means of other methods.

Figure 4.9: Conflicting routes.

The pseudocode for UC is much more detailed than the pseudocode of the other crossover operators, because it is considered necessary in order to explain it well enough. The pseudocode is as follows:

UCrossover(P1, P2)

NRP1 ← number of routes in P1
NRP2 ← number of routes in P2
offjind ← 0

for n ← 0 **to** NRP1
    for m ← 0 **to** NRP2
        if P1[n] conflicts with P2[m]
            $conflicts$[n][m] ← 1
    end for
end for

calculate $R_{P1,r}$ for all routes r in P1 and put the indices in *valueP1* according to a descending order of $R_{P1,r}$
calculate $R_{P2,r}$ for all routes r in P2 and put the indices in *valueP2* according to a descending order of $R_{P2,r}$

for s ← 0 **to** NRP1
    $possCandP1$[s] ← 1
end for
continues on next page

---

from previous page

for s ← 0 **to** NRP1
    possCandP1[s] ← 1
end for

P1ind ← 0
P2ind ← 0

while P1ind < NRP1 AND P2ind < NRP2 **do**
    foundP1 ← false
    foundP2 ← false

    while foundP1 = false AND P1ind < NRP1 **do**
        if *possCandP1*[*valueP1*[P1ind]] = 1
            offspring[offjind] ← P1[P1ind]
            for u ← 0 **to** Length(*conflicts*[0])
                if *conflicts*[*valueP1*[P1ind]][u] = 1
                  possCandp2[u] ← 0
            end for
            foundP1 ← true
            offjind ← offjind + 1
        end if

        P1ind ← P1ind + 1
    end while

    while foundP2 = false AND P2ind < NRP2 **do**
        if *possCandP2*[*valueP2*[P2ind]] = 1
            offspring[offjind] ← P2[P2ind]
            for v ← 0 **to** Length(*conflicts*)
                if *conflicts*[v][*valueP2*[P2ind]] = 1
                  possCandP1[v] ← 0
            end for
            foundP2 ← true
            offjind ← offjind + 1
        end if

        P2ind ← P2ind + 1
    end while
end while

put all other customers into *customersLeft*
sweep the customers in *customersLeft* into feasible routes and put into the *offspring*

return *offspring*

In the HLC the Sweep Algorithm is applied to generate feasible routes from the customers that are left out, see section 4.2.3. The Sweep Algorithm is also used here.

4.3   The Mutation Operator

The aim of the mutation operator is to make random changes to the offspring, which can not be done by a crossover operator. In chapter 2.2.6, the general idea behind the mutation operator is explained. One mutation operator was used in the project that is the Simple Random Mutation. The operator is discussed in the next subsection.

4.3.1   Simple Random Mutation

The *Simple Random Mutation* (SRM) performs a very simple mutation, which moves a single customer at a time within the same solution. It is not desirable to apply only the mutation so it can be assumed that a crossover operator has been applied before and it has generated an offspring, which SRM is applied to. There is a possibility to apply the crossover operators with a certain probability to slow down the convergence but there is no reason to do that when using a steady-state algorithm where only few solutions are changed within the population in each iteration. The pseudocode for SRM is:

Mutation(offspring, pSame)

NR ← number of routes in *offspring*
jMut ← random(0,NR)
RL ← number of customers in route jMut
kMut ← random(0,RL)
cust ← *offspring*[jMut][kMut]

delete *cust* from jMut

prob ← random(1,100)

if prob < pSame
        bestInd ← [jMut, BestInsertion(*cust*,jMut)]

else do j ← random(0, NR)
        while j = jMut
          bestInd ← [j, BestInsertion(*cust*,j)]

insertCust(*cust*, *bestInd*)

return *offspring*

The operator starts by randomly selecting a route jMut in the offspring and then a customer within the route is chosen randomly. The customer is deleted from the offspring to avoid duplication later on. It can either be inserted into the same route again or into another randomly chosen route within the same solution, controlled by the parameter *pSame*. The probability of being inserted into the same route again is equal to pSame. When the customer is inserted into the same route again, it is possible that it will be inserted in the same place as before and resulting in the same offspring again. The function BestInsertion returns the index indicating the best place for the customer to be inserted. It measures the pay off in the same way as in SRC, see section 4.2.1, but the insertion route has already been selected. The pay off of inserting customer k between customers c_m and c_{m+1} is:

$$p_m = cost(c_m, c_{m+1}) - cost(c_m, k) - cost(k, c_{m+1}) \tag{4.7}$$

where $cost(c_m, c_{m+1})$ is the cost of travelling between c_m and c_{m+1}.

It is important not to cause too much disruption using mutation, because that will ruin the characteristics of the population and result in a random search. The steady state algorithm only generates one offspring in each iteration, which the mutation is applied on with a certain probability called *mutation rate*, see chapter 2.2.2. Figure 4.10 illustrates very clearly how the result can get worse if mutation is applied too often. The results are obtained from using Simple Random Crossover, Repairing Operator, explained in next section, and Steepest Improvement Algorithm for four problems of different sizes; A-n32-k5, A-n44-k6, A-n60-k9 and A-n80-k10 from [20]. The numbers are average of five runs with a population size 50 and 10000 iterations. The mutation operator was applied 1, 2 and 3 times with pSame = 30%.

Figure 4.10: The effect on the performance of applying SRM number of times. When the number of application of SRM is increased the difference from optimum gets larger.

4.4   The Supporting Operators

The supporting operators are applied to a single solution similar to the mutation operators but they have a different purpose. They were developed in order to support the different crossover operators by repairing their flaws. Two operators were designed; Repairing Operator and Geographical Merge. They are explained in the following subsections.

4.4.1   Repairing Operator

In the development process of both SRC and BOC it seemed as if they often generated offsprings that were far from being infeasible, i.e. the violation of the capacity constraint was relatively large. That caused the penalty term of the fitness value to be too predominant. An experiment was made by limiting the size of the subroute to 1, 2 or 3 customers. If that would give better results, it was necessary to find a way to reduce the violation of the capacity. The graphs in figure 4.11 illustrate the effect of reducing the subroute length. They show the percentage difference from optimum when the two crossover operators as well as Simple Random Mutation (pSame = 30% and rate = 100%) and Steepest Improvement Algorithm were used to solve four problems of different sizes; A-n32-k5, A-n44-k6, A-n60-k9 and A-n80-k10 from [20]. The results are obtained with 5 runs, a population of 50 individuals and 10000 iterations. The numerical results are given in appendix B.

There is a clear tendency in both graphs that the performance improves when the number

Figure 4.11: The performance of the Simple Random Crossover and Biggest Overlap Crossover using different subroute sizes. Max 3 indicates that the subroute contains maximum 3 customers, etc.

of customers in the subroute is limited. The difference is larger for Simple Random Crossover, which most likely can be explained with Biggest Overlap Crossover taking the capacity into consideration when it inserts the subroute into a route. The Simple Random Crossover chooses the insertion route by random, so often the subroute will be inserted into a route with high total demand.

Too predominant penalty term can either be caused by the penalty formula or because the offsprings are in general relatively far from attaining the capacity constraint. Even though it is the intention to allow infeasible solutions at the early iterations, the population can have difficulties in improving if many generated offsprings have too large total demands. Therefore, it was decided to try to make an operator, which could bring the offspring a bit closer to feasibility. *Repairing Operator* (RO) was designed to carefully shorten the routes, starting with the routes violating the capacity constraint at most. It selects the route with the largest total demand and if it does not attain the capacity constraint, a customer is randomly chosen and moved to the route with the smallest total demand.

The pseudocode is as follows:

Repair(offspring)

jMax ← the index of the route in the *offspring* with largest total demand
maxDem ← the total demand of route jMax
jMin ← the index of the route in the *offspring* with smallest total demand

if maxDem > K
      randomly choose a customer from route jMax
      delete the customer from route jMin
      insert the customer at the end of route jMin

return *offspring*

4.4.2 Geographical Merge

The role of *Geographical Merge* (GM) is to merge two routes, which are close to each other, within the same individual. Horizontal Line Crossover and Uniform Crossover, explained in section 4.2.3 and 4.2.4 respectively, tend to generate offsprings with relatively short routes. In order to illustrate that, the average difference from the capacity was calculated, separately for the routes with too much total demand and for those with total demand within the capacity of the vehicles. Equation 4.8 illustrates the formula, which was used to calculate the difference d_p for routes attaining the capacity constraint, where p stands for plus. If a route does no contain any routes of either kind the difference is set to zero. The formula for the routes violating the capacity constraint is corresponding. The following formula shows the difference for the routes attaining the capacity constraint.

$$d_p = \frac{\sum(K - totdem_p)}{nr_{p,r}} \tag{4.8}$$

where:

K	is the capacity of the vehicles,
$totdem_p$	is the total demand of a route r where it is less than K and
nr_p	denotes number of routes with total demand less than K.

In the following figures the results, which were obtained by solving problem instance A-n80-k10 from [20] using the two crossover operators, Simple Random Mutation (pSame = 30% and rate = 100%) and Steepest Improvement Algorithm. The algorithm was run once with a population size 50 and 10000 iterations. The capacity of the vehicles is 100.

Figure 4.12: Average difference from capacity using Horizontal Line Crossover.

Figure 4.13: Average difference from capacity using Uniform Crossover.

The positive difference in the figures indicates how much more demand could be added to the routes on average in order to fill the vehicles and the negative difference shows how much more demand than can be carried by the vehicles is in the routes on average. Particularly when using HLC, there is a relatively large gap between the total demand and the capacity of the vehicles for feasible routes. The gap is at least 10 capacity

units, besides from the first iterations. The gap is smaller for UC or usually more than 5 capacity units. Furthermore, figure 4.13 clearly shows that UC does seldom generate infeasible routes and those that are generated are usually very close to being feasible.

The UC operator uses both the capacity of the routes in addition to the distance between them to combine them. The distance between every two routes is measured using their bounding boxes, see section 4.2.2. The pseudocode is rather detailed in order to explain the operator well enough.

Merge(offspring)

NR ← number of routes in *offspring*

for g ← 0 **to** NR
    dem1 ← demand of route *offspring*[g]
    for h ← g **to** NR
        dem2 ← demand of route *offspring*[h]
        if dem1 + dem2 ≤ K
          goodPairs[g][h] ← 1
    end for
end for

biggestoverlap ← 0
closest ← MAX

for e ← 0 **to** NR
    for f ← e **to** NR
        if e ≠ f AND *goodPairs*[e][f] = 1
          bol ← findOverlap(e,f)
          if bol ≠ 0 AND bol > biggestoverlap
            biggestoverlap ← bol
            olroute1 ← e
            olroute2 ← f
        else if biggestoverlap = 0
          cl ← findClosest(e,f)
          if cl < closest
            closest ← cl
            clroute1 ← e
            clroute2 ← f
    end for
end for
 insert route clroute2 at the end of clroute1

return *offspring*

The operator is applied on a single individual. It starts by identifying pairs of routes that together have the total demand within the limits of the capacity, thus it does not generate infeasible routes. Then the pair of routes closest to each other is selected. If the bounding boxes of the routes overlap, the pair with the largest overlap is chosen. Otherwise, it selects the pair with the shortest distance between the bounding boxes. Finally, the routes are merged together by adding the one route at the end of the other.

4.5   Summary

In this chapter the development process of the fitness value and 7 operators is discussed. Four crossover operators are explained; Simple Random Crossover, Biggest Overlap Crossover, Horizontal Line Crossover and Uniform Crossover. Simple Random Mutation, the only mutation operator, is described and also two supporting operators; Repairing Operator and Geographical Merge. In the next chapter some implementation issues are discussed.

Chapter 5

Implementation

In this chapter implementation issues are discussed, such as information about the computer, the experiments that were carried out and computation of Euclidean distances, among other things.

Computational experiments were carried out on a Zitech PC with a 800 MHz AMD Athlon(tm) processor. The program was implemented using the Java 1.3.1_01 Standard Edition and compiled using the java compiler.

The CPU-time was calculated using the function *currentTimeMillis()* in Java. It returns the current time in milliseconds since midnight January 1, 1970 GMT. The function is used right before the algorithm is applied in the code and again afterwards. The execution time of the algorithm is the difference between two measurements.

The cost of travelling between every two customers was calculated using Euclidean distances. Some problem instances have been calculated using real numbers in which the best results are presented. Others are calculated with integers. When integers are used, the rounding function *Nearest Integer* function *nint* was used to convert the Euclidean distances into integers. The function returns the integer closest to the given number and half-integers are rounded up to the smallest integer bigger than the number, i.e. $nint(1.4) = 1$, $nint(3.8) = 4$ and $nint(5.5) = 6$. When Java rounds real numbers, all numbers are rounded down to the nearest integer. Therefore, a constant 0.5 was added to all the real numbers in order to simulate nint. For instance, Java rounds 3.8 down to 3 by Java but by adding 0.5 to it, $3.8 + 0.5 = 4.3$, it is instead rounded "up" to 4!

In section 2.2.2 the representation of the VRP solution used in this project is described. The population is implemented as a 3-dimensional array M $\times$ dim $\times$ dim, where M is the population size and dim is the number of customers in each problem. So Pop[2][5][3] = 5 means that customer no. 5 is visited number 4 in route 6 in the third individual in the population (indices start in 0 in Java).

The advantage of presenting the population in this way is that it is very simple to implement. Furthermore, it simplifies communication and gives a good overview. The drawback is that it contains a lot of empty spaces that use memory, where either the number of customers is less than dim or the number of routes is less than dim. The memory need

of the population could be reduced by using upper bound on the number of customers within a route or on the number of routes, based on the problem to be solved. However, one should be careful about using upper bounds because Java is not able to enlarge an array dynamically. Thus if the program tries to insert a customer outside the array the program terminates.

In order to save computational time, the total cost of each individual was kept in separate arrays for the cost and the penalty. Since the program used a steady-state algorithm (see section 2.2.2) only few solutions were changed in each iteration. When a solution was changed, it was noted in the cost and the penalty array and next time it was selected the cost was recalculated. In a similar way the bounding boxes were also kept in an array.

Chapter 6

Parameter Tuning

Many parameters are involved in solving VRP with GA. The values of those need to be tuned. Good parameter tuning is important for all GA to generate good solutions. Usually, it is complex to tune the parameters because the relationship between them can be rather complicated and unclear. Most algorithms contain many parameters and the more they get the more time consuming the tuning becomes. Ideally, all parameters should be tuned but that is in general too time consuming. However, it is important to tune as many as possible. In the following section all parameters are listed, it is explained how they were tuned and the values they were given are shown.

6.1   The Parameters and the Tuning Description

The parameters are divided into two groups of more or less important parameters. In the group of more important parameters are those that are believed to have more influence on the effectivity of the algorithm. In the other group are parameters of less significance and the effect of some of them has been discussed in the previous chapter. The parameters are listed in the two tables below:

1.	The Population size (M).
2.	The rate of SRM.
3.	The rate of RO.
4.	The rate of GM.

Table 6.1: The more important parameters.

5.	A scaling factor for α in the penalty function.
6.	The power in the penalty function.
7.	The power in the α function.
8.	The tournament size when choosing parent solutions.
9.	The probability of Tournament Selection when choosing parent. solutions.
10.	The tournament size when choosing an individual leaving the population.
11.	The probability of Tournament Selection when choosing an individual out of population.
12.	The rate of SRC.
13.	The subroute length in SRC.
14.	The rate of BOC.
15.	The subroute length in BOC.
16.	Number of routes closest to the subroute in BOC.
17.	The rate of HLC.
18.	The rate of UC.
19.	The scaling factor for evaluating the quality of the routes in UC.
20.	The probability of inserting the subroute into the same route again in SRM (pSame).
21.	Number of customers moved in SRM.
22.	Number of customers moved in RO.

Table 6.2: The less important parameters.

Although it is desirable to tune all possible parameters, it is hardly never possible for a program of this size because it takes too much time. Here only the more important parameters are tuned. The reasons for why the less important parameters are considered so and why they are not tuned are the following:

Parameter 5: The effect of different values of α was discussed shortly in section 4.1 and because of limited time it is not tuned here.

Parameters 6 and 7: The power in neither the penalty function nor the α-function are tuned, because that will have similar effect as tuning of α. For simplification and to save time they are kept constant.

Parameters 8 and 10: The tournament sizes of neither selection of parent solutions nor the solution leaving the population are tuned. The convergence depends both on the selection pressure and the population size and to save time the convergence will be balanced by tuning the population size with constant tournament size of two.

Parameters 9 and 11: It is not important to tune the probability of tournament selection when using a steady-state algorithm. On the other hand, it can be necessary

to reduce the selection pressure by applying the tournament selection with a certain probability when generational algorithms are used, see page 20.

Parameters 12, 14, 17 and 18: The probability of applying a crossover operator is not tuned. Skipping a crossover in some iterations can be done to be better able to preserve a part of the population from one iteration to another. That is not important here because a steady-state algorithm is used.

Parameters 13 and 15: The lengths of the subroutes in SRC and BOC are not tuned because it is important to keep some randomness in the crossover operators.

Parameters 16, 19 and 20: The number of routes or candidates in BOC, is not tuned since it will most likely not have much influence compared to the time it takes to tune it. For the same reasons a possible scaling factor in UC and pSame are not tuned. The parameter pSame is set to 30%, which has worked out fine during the development process of the project.

Parameters 21 and 22: Instead of tuning the number of customers moved in SRM and RO directly they are tuned at some degree by tuning the rate of the operators where it is possible to apply them twice.

Two algorithms are defined; a fast algorithm with **10000** iterations and a slow algorithm with **100000** iterations. The algorithms are tuned for different sets of population sizes but parameters 2, 3 and 4 are tuned for the same values for both algorithms. Since there is not a linear correlation between the population size and the number of iterations, the slow algorithm is not necessarily expected to perform better than the fast algorithm.

The problem instances that are used for both tuning and testing are divided into two groups of small instances with less than 100 customers and large instances with 100 customers or more. The size of the large instances limits the possible population size to 100 due to the memory of the computer that is used. Furthermore, running the slow algorithm using a relatively large population for large problems is extremely time consuming. Thus, the large instances are only tuned for the fast algorithm and smaller populations. The values of parameters 2, 3 and 4 for the slow algorithm and the large instances are borrowed from the tuning of the small instances for the slow algorithms. The possible values of parameters are chosen based on the experience that has been gained during the development process. The following tables show the possible values for small problems using fast and slow algorithm and large problems using fast algorithm.

	Parameters	Fast	Slow
1	Population size (M)	50,100,200	200,400
2	SRM rate (%)	0,50,100,200	0,50,100,200
3	RO rate (%)	0,50,100,200	0,50,100,200
4	GM rate (%)	0,50,100,200	0,50,100,200

Table 6.3: Possible values of parameters for small problems using a fast or a slow algorithm.

	Parameters	Fast
1	Population size (M)	50,100
2	SRM rate (%)	0,50,100,200
3	RO rate (%)	0,50,100,200
4	GM rate (%)	0,50,100,200

Table 6.4: Possible values of parameters for large problems using fast algorithm.

Ideally, it is preferred to tune all possible combinations of operators. For time reasons, that will not be done here. Instead 8 combinations of operators are considered. Each crossover operator is tuned with SRM and a corresponding supporting operator, RO or GM. Furthermore, they are tuned with and without SIA. The different combinations are:

1. SRC, SRM, RO and SIA
2. SRC, SRM and RO
3. BOC, SRM, RO and SIA
4. BOC, SRM and RO
5. HLC, SRM, GM and SIA
6. HLC, SRM and GM
7. UC, SRM, GM and SIA
8. UC, SRM and GM

It is essential not to use the same problem instances for tuning and testing. Four problem instances are chosen for the parameter tuning; two small problems and two large problems. The final testing is performed with 12 small problem instances and 7 large problem instances, thus the sets of tuning problems are $\frac{1}{7}$ and $\frac{2}{9}$ of the whole set of problems. Table 6.5 illustrates the problems. The bold results correspond to optimal values and the regular results denote the best know values.

Small			Large		
Problem	Size	**Optimal**/ Best known value	Problem	Size	**Optimal**/ Best known value
A-n32-k5	32	**784**	P-n101-k4	101	**681**
vrpnc1	50	**524,61**	vrpnc12	100	819,56

Table 6.5: Small and large tuning problems, their sizes and best known or optimal values.

The tuning instances are chosen among the different types of problems that are used in the final testing. Problems A-n32-k5 and P-n101-k4 belong to sets A and P of the Augerat et al. problems in [20] and problems vrpnc1 and vrpnc12 belong to the Christofides, Mingozzi and Toth (CMT) problems in [17]. They are often used in the literature to compare different metaheuristics when solving VRP.

The optimal values of the Augerat et al. problems are given in integers and the best known values of the CMT problems are given in real numbers. For simplification all tuning is performed using integers to compute the cost. The results are presented as percentage difference from the optimal or best known value. Using integers does not affect the tuning results because the best suited parameters are chosen using the relative difference in performance between the different combinations of the parameters. However, when an algorithm is close to finding the optimal or best known values of the problems it is possible that the results will so sometimes be negative. That can be explained by the accumulated error of using integers instead of real numbers.

6.2   The Results of Tuning

The results of the parameter tuning are revealed in tables in appendix C, which show the percentage difference from optimum and the standard deviation. When the parameters were chosen, an effort was made to find the best combination for both problem instances. There was in general a fine consistency in the results of both tuning problems for all combinations of operators. Tables 6.6 to 6.9 illustrate the tuning results. The combinations of operators are illustrated on page 68.

	Combination 1			Combination 2		
Type	M	SRM (%)	RO (%)	M	SRM (%)	RO (%)
Small and fast	100	50	200	200	0	200
Small and slow	400	50	200	400	100	100
Large and fast	50	0	200	50	0	200
Large and slow	100	50	200	100	100	100

Table 6.6: Tuned parameters for combinations 1 and 2, i.e. SRC with and without SIA.

	Combination 3			Combination 4		
Type	M	SRM (%)	RO (%)	M	SRM (%)	RO (%)
Small and fast	50	100	100	50	50	100
Small and slow	200	50	100	400	100	100
Large and fast	50	50	200	50	100	100
Large and slow	100	50	100	100	100	100

Table 6.7: Tuned parameters for combinations 3 and 4, i.e. BOC with and without SIA.

	Combination 5			Combination 6		
Type	M	SRM (%)	GM (%)	M	SRM (%)	GM (%)
Small and fast	50	50	50	50	100	50
Small and slow	200	50	50	200	50	0
Large and fast	50	100	50	50	100	200
Large and slow	100	50	50	100	50	0

Table 6.8: Tuned parameters for combinations 5 and 6, i.e. HLC with and without SIA.

	Combination 7			Combination 8		
Type	M	SRM (%)	GM (%)	M	SRM (%)	GM (%)
Small and fast	100	100	100	200	100	50
Small and slow	400	50	50	200	100	200
Large and fast	100	100	50	50	100	100
Large and slow	100	50	50	100	100	200

Table 6.9: Tuned parameters for combinations 7 and 8, i.e. UC with and without SIA.

6.3 Summary

All possible parameters of the program have been listed. The more important parameters have been tuned for 8 different combinations of operators and their values are given in tables 6.6 to 6.9. In the next chapter the final testing is described and the results are illustrated.

Chapter 7

Testing

In this chapter the testing of the operators is described and the results are illustrated. The operators are tested by the the means of 8 different combinations that are illustrated on page 68 and the testing is performed using the parameters found in tables 6.6 to 6.9 in the previous chapter. The tests are made in two steps. Firstly, all combinations of operators are tested in order to choose the best combination. Secondly, the best combination is applied to an additional number of problems and the results are compared to recently proposed results of other metaheuristics.

7.1   The Benchmark Problems

The testing is performed on 12 small and 7 large problems, which are listed in the tables below. It would have been desirable to perform the testing using more large problems but large problems with known optimal or best values were difficult to find. The optimal values are illustrated bold.

Problem	Size	Optimal/ Best known values
E-n33-k4	33	**835**
B-n41-k6	41	**829**
F-n45-k4	45	**724**
P-n50-k7	50	**554**
A-n60-k9	60	**1354**
B-n68-k9	68	**1272**
F-n72-k4	72	**237**
E-n76-k7	76	**682**
P-n76-k5	76	**627**
B-n78-k10	78	**1221**
A-n80-k10	80	**1763**
vrpnc2	75	835,26

Table 7.1: The sizes and the optimal or best known values of the small problems.

Problem	Size	Optimal/ Best known values
F-n135-k7	135	**1162**
E-101-k8	101	817
G-n262-k25	262	6119
vrpnc3	100	826,14
vrpnc4	150	1028,42
vrpnc5	199	1291,45
vrpnc11	120	1042,11

Table 7.2: The sizes and the optimal or best known values of the large problems.

The problem instances belong to different groups of known problems. Problems E-n33-k4, E-n76-k7 and E-n101-k8 belong to the Christofides and Eilon problems. B-n41-k6, P-n50-k7, A-n60-k9, B-n68-k9, P-n76-k5, B-n78-k10 and A-n80-k10 belong to the problem sets A, B and P of Augerat et al. Problems F-n45-k4, F-n72-k4 and F-n135-k7 belong to the group of problems proposed by Fisher. G-n262-k25 is a Gillet and Johnson problem. All these problems are from [20]. The vrpnc-problems of Christofides, Mingozzi and Toth (CMT) are the problems in which most recently proposed results of metaheuristics for VRP are presented. They are from [17].

7.2   Test Description

Each crossover operator is tested with and without SIA, i.e. for 8 different combinations of operators. Each combination is tested four times, i.e. for both small and large problems and both fast and slow algorithm.

The first part of the testing is performed in order to choose the best possible combination for each problem size and each type of algorithm. Three small and three large problem instances are selected from tables **??** and 7.2: B-n41-k6, B-n68-k9 and vrpnc2 for small problems and vrpnc3, vrpnc9 and vrpnc11 for large problems. The different combinations are applied to these problems and the best one is selected, primarily considering the accuracy. That is measured by the difference from the optimum or the best known value of each problem. The selection is based on the average performance of **10** runs. Boxplots help with illustrating the difference in performance between the operators. Consider the results in a vector v_c where c is the index of the considered combination. A boxplot draws the median of v_c, the minimum and the maximum value, the 1^{st} and the 3^{rd} quartile.

In the second part of the testing, the best combination is applied to the rest of the problems within each size for each type of algorithm. For each size the results of the fast and the slow algorithm are compared. The more accurate algorithm is chosen. The results are presented for solving the small CMT problems with the best algorithm for the small problems and the large CMT problems with the best algorithm for the large problems. The results are compared to the proposed results of three Tabu Search heuristics; Taburoute, Taillard's Algorithm and Adaptive memory [4], and the Hybrid Genetic Algorithm 2.1.4.

7.3   The Results

In the following 7 subsections the results are presented. The first three subsections involve the testing for the small problems. The next three sections illustrate the results for the large problems. In the final subsection, the results of the comparison with proposed the Tabu Search heuristics and the Hybrid Genetic Algorithm are stated.

7.3.1   Small Problems and Fast Algorithm

In order to choose the best combination of operators, the combinations are compared using the fast algorithm to solve the small problems. The results of solving the three small problem instances are illustrated in the following four tables. Each table illustrates the results of using one of the crossover operators with or without SIA. The results are presented as the average percentage difference from optimum or best know value, standard deviation and time. The combinations are explained on page 68 and the parameters are shown in tables 6.6 to 6.9.

Problem	Combination 1			Combination 2		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	5,1	3,44	5120	7,55	4,31	2273
B-n68-k9	14,11	3,8	10297	23,23	5,83	5094
vrpnc2	16,21	2,33	10453	30,96	6,48	6305
Average	11,81	3,19	8623	20,58	5,54	4557

Table 7.3: Results of solving the small problems with the fast algorithm using SRC with and without SIA.

Problem	Combination 3			Combination 4		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	8,6	4,81	4291	14,68	4,44	1999
B-n68-k9	20,74	3,66	9371	39,16	6,46	4666
vrpnc2	14,13	2,67	13220	32,4	3,59	5894
Average	14,49	3,71	8961	28,75	4,83	4186

Table 7.4: Results of solving the small problems with the fast algorithm using BOC with and without SIA.

Problem	Combination 5			Combination 6		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	12,34	4,99	2768	16,61	6,35	2545
B-n68-k9	11,47	2,47	5885	20,47	4,11	5771
vrpnc2	25,41	4,4	9111	49,03	5,8	7206
Average	16,41	3,95	5921	28,70	5,42	5174

Table 7.5: Results of solving the small problems with the fast algorithm using HLC with and without SIA.

Problem	Combination 7			Combination 8		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	8,06	4,23	5008	15,96	4,14	2685
B-n68-k9	12,51	4,25	11027	19,69	3,69	6084
vrpnc2	13,26	1,96	17548	23,49	2,15	7286
Average	11,28	3,48	11194	19,71	3,32	5342

Table 7.6: Results of solving the small problems with the fast algorithm using UC with and without SIA.

The boxplot below illustrates the results graphically. Note that the performance of each combination is presented as average in the tables and as the median in the boxplot. Also note that combinations 1, 3, 5 and 7 are with SIA and 2, 4, 6 and 8 are without SIA.

Figure 7.1: The performance of the different combinations for the fast algorithm and the small problems. Combination 7 gives the best results.

The results illustrate that **combination 7** with UC and SIA performs best. Although combination 1 with SRC and SIA provides almost as good results on average. It is quite clear that SIA has rather great improving effect for all the crossover operators but it takes time. For instance, SIA reduces the difference from optimum by almost half for BOC, see table 7.4. BOC and HLC do not perform well enough. Combination 7 is the most time consuming one combination and combination 5 is the least time consuming of those that are supplied with SIA.

Combination 7 is now applied to the rest of the small problems using the fast algorithm and the results of all problems are illustrated in the table 7.7. The results of B-n41-k6, B-n68-k9 and vrpnc2 are obtained from table 7.6

Problem	Diff. from opt. (%)	Std. dev. σ	Time (ms)
E-n33-k4	4,14	1,99	2857
B-n41-k6	8,06	4,23	5008
F-n45-k4	4,91	2,59	753
P-n50-k7	7,27	3,64	4601
B-n56-k7	7,42	2,38	6019
A-n60-k9	12,82	1,38	6078
B-n68-k9	12,51	4,25	11027
F-n72-k4	25,3	7,62	14101
P-n76-k5	10,25	2,58	12075
B-n78-k10	13,58	1,91	10260
A-n80-k10	16,91	3,5	10836
vrpnc2	13,26	1,96	17548
Average	11,37	3,17	8430

Table 7.7: Results of solving all the small problems with the fast algorithm using combination 7, i.e. UC and SIA.

7.3.2   Small Problems and Slow Algorithm

In this subsection the combinations of operators are compared for solving the small problems with the slow algorithm. The results are shown in the tables 7.8 to 7.11. Each table includes the results of each crossover operator with or without SIA. Different combinations are explained on page 68 and the tuned parameters are shown in tables 6.6 and 6.9.

Problem	Combination 1			Combination 2		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	2,41	1,73	46842	4,91	2	20774
B-n68-k9	4,03	1,8	100492	14,25	5,56	46985
vrpnc2	6,08	0,94	96782	18,41	1,66	59957
Average	4,17	1,49	81372	12,52	3,07	42572

Table 7.8: Results of solving the small problems with the slow algorithm using SRC and SIA.

Problem	Combination 3			Combination 4		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	5,05	3,31	42002	6,12	3,56	19865
B-n68-k9	5,86	3,4	93619	19,51	6,11	45559
vrpnc2	5,28	2,09	142070	18,57	2,95	55418
Average	5,40	2,93	92564	14,73	4,21	40281

Table 7.9: Results of solving the small problems with the slow algorithm using BOC with and without SIA.

Problem	Combination 5			Combination 6		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	5,17	2,05	42782	10,75	3,87	20593
B-n68-k9	9,24	0,99	83666	12,54	2,18	44186
vrpnc2	31,81	3,74	68152	26,58	3,24	59867
Average	15,41	2,26	64867	16,62	3,10	41549

Table 7.10: Results of solving the small problems with the slow algorithm using HLC with and without SIA.

Problem	Combination 7			Combination 8		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
B-n41-k6	0,83	0,57	54246	3,63	2,83	37645
B-n68-k9	6,58	2,27	101244	10,53	1,89	85522
vrpnc2	7,2	2,07	165125	11,99	2,4	113309
Average	4,87	1,64	106872	8,72	2,37	78825

Table 7.11: Results of solving the small problems with the slow algorithm using UC with and without SIA.

The results are also illustrated in the boxplot below. Note that the performance of each combination is presented as average in the tables and as the median in the boxplot. Also note that combinations 1, 3, 5 and 7 are with SIA and 2, 4, 6 and 8 are without SIA.

Figure 7.2: The performance of the different combinations for the slow algorithm and the small problems. Combination 1 gives the best results, but combinations 3 and 7 also perform quite well.

Combination 1 with SRC and SIA is the best combination. However, combinations 3 with BOC and SIA and 7 with UC and SIA do also perform well. As for the fast algorithm, application of SIA improves the results of the crossover operators considerably. On the other hand, combinations with SIA are more time consuming. Both combinations with HLC perform particularly bad compared to the other combinations.

In table 7.12 the results of applying combination 1 using the slow algorithm to solve the small problems are shown. The results of problems B-n41-k6, B-n68-k9 and vrpnc2 are from table 7.8.

Problem	Diff. from opt. (%)	Std. dev. σ	Time (ms)
E-n33-k4	0,6	0,7	31216
B-n41-k6	2,41	1,73	46842
F-n45-k4	0,91	0,23	61341
P-n50-k7	2,09	1,1	42566
B-n56-k7	1,37	0,18	53058
A-n60-k9	3,29	1,52	55650
B-n68-k9	4,03	1,8	100492
F-n72-k4	8,86	1,24	156091
P-n76-k5	5,56	2,87	110854
B-n78-k10	6,67	1,27	86476
A-n80-k10	8,09	1,1	93596
vrpnc2	6,08	0,94	96782
Average	4,16	1,22	77914

Table 7.12: Results of solving all the small problems with the slow algorithm using combination 1, i.e. SRC and SIA.

7.3.3 Comparison of Fast and Slow Algorithm for Small Problems.

Combination 7 with UC and SIA performed best for the fast algorithm and combination 1 with SRC and SIA provided the best results for the slow algorithm. By comparing the results of tables 7.7 and 7.12 respectively, **the slow algorithm with combination 1** turns out to give the best results for the small problems. The difference in the performance is roughly 7%. Naturally, the slow algorithm is more time consuming. However, by looking closer at the average time it is observed that 77914 ms are a little bit less than 1,3 min. That is not particularly bad compared to both Taburoute and Hybrid Genetic Algorithm in table 7.23 on page 85. Thus, the slow algorithm is selected and it will be used to compare the algorithm to other metaheuristics.

7.3.4 Large Problems and Fast Algorithm

In order to choose the best combination of the operators, the combinations are compared using the fast algorithm to solve the large problems. The following four tables illustrate the results. Each table presents the average percentage difference from optimum, standard deviation and time of using one of the crossover operators with and without SIA. The combinations are explained on page 68 and the corresponding parameters are illustrated in tables 6.6 to 6.9 in the previous chapter.

Problem	Combination 1			Combination 2		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	15,84	3,13	29750	32,69	7,65	14954
vrpnc4	30,22	4,13	53300	75,43	7,89	33057
vrpnc11	51,91	7,61	54822	62,98	27,88	21717
Average	32,66	4,96	45957	57,03	14,47	23243

Table 7.13: Results of solving the large problems with the fast algorithm using SRC with and without SIA.

Problem	Combination 3			Combination 4		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	18,98	1,27	20358	49,31	4,85	14253
vrpnc4	42,04	2,42	36679	101,19	6,66	31769
vrpnc11	51,64	7,15	33359	71,35	4,82	20769
Average	37,55	3,61	30132	73,95	5,44	22264

Table 7.14: Results of solving the large problems with the fast algorithm using BOC with and without SIA.

Problem	Combination 5			Combination 6		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	8,89	1,93	28251	35,86	4,97	8163
vrpnc4	21,96	2,54	45412	85,93	16,76	17885
vrpnc11	28,93	2,32	69858	121,92	14,53	10596
Average	19,93	2,26	47840	81,24	12,09	12215

Table 7.15: Results of solving the large problems with the fast algorithm using HLC with and without SIA.

Problem	Combination 7			Combination 8		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	8,69	1,23	35861	26,33	2,75	11565
vrpnc4	15,47	1,99	55290	41,04	4,47	26519
vrpnc11	24,93	3,3	53629	50,93	8,3	16606
Average	16,36	2,17	48260	39,43	5,17	18230

Table 7.16: Results of solving the large problems with the fast algorithm using UC with and without SIA.

The results are also illustrated in the boxplot below. Note that the performance of each combination is presented as average in the tables and as the median in the boxplot. Also note that combinations 1, 3, 5 and 7 are with SIA and 2, 4, 6 and 8 are without SIA.

Figure 7.3: The performance of the different combinations for the fast algorithm and the large problems. Combination 7 gives the best results, but combination 5 also performs quite well.

Combination 7 with UC and SIA is clearly the best one here. Combination 5 with HLC and SIA performs a bit worse. Combinations 1 with SRC and SIA and 3 with BOC and SIA do not provide good enough results. The performance of all combinations without SIA is unacceptable, even though they are less time consuming. Combinations 1, 5 and 7 are almost equally time consuming but combination 3 is quite faster.

The following table illustrates the results of using the fast algorithm with combination 7 to solve all the large problems:

Problem	Diff. from opt. (%)	Std. dev. σ	Time (ms)
E-101-k8	8,92	3,29	25371
F-n135-k7	28,15	4,48	57096
G-n262-k25	14,17	1,16	131852
vrpnc3	8,69	1,23	35861
vrpnc4	15,47	1,99	55290
vrpnc5	19,84	1,8	93021
vrpnc11	24,93	3,3	53629
Average	17,17	2,46	64589

Table 7.17: Results of solving large problems with the fast algorithm using combination 7, i.e. UC and SIA.

7.3.5   Large Problems and Slow Algorithm

The results of solving the large problems using the slow algorithm and different combinations of operators are illustrated in the following four tables. The average performance is presented as difference from optimum or best known values, standard deviation and time. The combinations are explained on page 68 and the parameters that are used are shown in tables 6.6 and 6.9 in the previous chapter.

Problem	Combination 1			Combination 2		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	6,00	2,3	200567	16,51	5,81	99297
vrpnc4	14,36	2,14	366507	37,07	3,69	125445
vrpnc11	31,37	7,32	386998	34,19	10,23	145847
Average	17,24	3,92	318024	29,26	6,58	123530

Table 7.18: Results of solving the large problems with the slow algorithm using SRC with and without and SIA.

Problem	Combination 3			Combination 4		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	10,87	3,64	185866	17,37	4,46	137669
vrpnc4	19,96	4,53	344188	42,37	4,29	346149
vrpnc11	28,57	5,17	312109	38,89	9,96	201582
Average	19,80	4,45	280721	32,88	6,24	228467

Table 7.19: Results of solving the large problems with the slow algorithm using BOC with and without SIA.

Problem	Combination 5			Combination 6		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	3,09	1,28	239140	21,85	2,84	95496
vrpnc4	15,81	2,96	457284	43,98	3,63	204748
vrpnc11	24,66	2,93	633075	53,75	14,75	130580
Average	14,52	2,39	443166	39,86	7,07	143608

Table 7.20: Results of solving the large problems with the slow algorithm using HLC with and without SIA.

Problem	Combination 7			Combination 8		
	Diff. from opt./best (%)	Std. dev. σ	Time (ms)	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
vrpnc3	3,27	1,78	202719	12,87	3,18	162174
vrpnc4	10,21	0,82	362012	21,11	3,46	270861
vrpnc11	18,63	1,04	374695	35,44	5,89	229125
Average	10,70	1,21	313142	23,14	4,18	220720

Table 7.21: Results of solving the large problems with the slow algorithm using UC with and without SIA.

The boxplot below illustrates the results. Note that the performance of each combination is presented as average in the tables and as the median in the boxplot. Also note that combinations 1, 3, 5 and 7 are with SIA and 2, 4, 6 and 8 are without SIA.

Figure 7.4: The performance of the different combinations for the slow algorithm and the large problems. Combination 7 gives the best results.

Combination 7 with UC and SIA performs best when using the slow algorithm to solve the large problems. All combinations with SIA perform relatively. Combinations 2, 4, 6 and 8 do not provide good enough results. Combination 3 with BOC and SIA is the least time consuming of those with SIA. Combinations 1 and 7 are almost equally time consuming and combination 5 requires the most time.

The following table shows the results of using the slow algorithm with combination 7 to solve all the large problems.

Problem	Diff. from opt./best (%)	Std. dev. σ	Time (ms)
E-101-k8	4,91	2,55	226514
F-n135-k7	17,98	3,36	537567
G-n262-k25	9,13	1,28	1308962
vrpnc3	3,27	1,78	202719
vrpnc4	10,21	0,82	362012
vrpnc5	14,28	1,73	889969
vrpnc11	18,63	1,04	374695
Average	11,20	1,79	557491

Table 7.22: Results of solving all the large problems with the slow algorithm using combination 7, i.e. UC and SIA.

7.3.6 Comparison of Fast and Slow Algorithm for Large Problems.

Combination 7 with UC and SIA performed best for both the fast algorithm and the slow algorithm. A comparison of the results of tables 7.17 and 7.22, shows that the slow algorithm performs better than the fast one. The difference in the performance is a little bit less than 6%. The slow algorithm is more time consuming. On average it requires 557491 ms, which accounts for approximately 9,3 min. As before, accuracy is considered more important than computation time, as long as it is within reasonable limits. Consequently, **the slow algorithm with combination 7** is chosen for the large problems.

7.3.7 Comparison of the Algorithm and other Metaheuristics

Now it is time to compare the best algorithms that have been obtained in this project to proposed results of the best metaheuristics. The slow algorithm with combination 1 performed best for the small problems and the slow algorithm with combination 7 provided the best results for the large problems. Hereafter, the algorithms will be referred to as *SRC-GA* and *UC-GA*, respectively.

In table 7.23 the final results are shown for 7 CMT problems. It was not the intention of this project to focus specially on the CMT problems but the results of the best metaheuristics are proposed by the means of these problems. The CMT problems are actually 14. Those that are not in the following table also include maximum route time and a drop time for each customer. The algorithm of this project does not support these factors and those problems are therefore eliminated.

The results of the Tabu Search heuristics Taburoute, Taillard and Adaptive memory are from [4]. They are explained shortly in section 2.1.4. The results of Taillard and Adaptive memory present the best performance of several runs. Information about the how the tests were made for Taburoute are missing. Although a comparison of the best performance of several runs and the average performance of ten runs is really unfair, it gives a clue about

the efficiency of the algorithms of this project. Luckily, Berger and Barkaoui present in [2] the result of the Hybrid Genetic Algorithm (HGA-VRP) as an average performance. Although, the number of runs is not given. The algorithm is explained shortly in section 2.1.4.

The results of SRC-GA for the small problems are taken from table 7.12 and the results of UC-GA for the large problems are taken from table 7.22. Problems vrpnc1 and vrpnc12 were used for parameter tuning. Usually, it is not preferable to present the results of the problems that were used for tuning. However, this is done here in order to make the comparison based on more problems. Problem vrpnc1 was recalculated using SRC-GA and the results of vrpnc12 were obtained using UC-GA.

	Taburoute		Taillard	Adaptive memory	HGA-VRP		SRC-GA and UC-GA	
Problem	Perf. (%)	Time (min)	Perf. (%)	Perf. (%)	Perf. (%)	Time (min)	Perf. (%)	Time (min)
vrpnc1	0,00	6,0	0,00	0,00	0,00	2,00	2,12	1,69
vrpnc2	0,06	53,8	0,00	0,00	0,57	14,33	6,08	1,61
vrpnc3	0,04	18,4	0,00	0,00	0,47	27,90	3,27	3,38
vrpnc4	0,08	58,8	0,00	0,00	1,63	48,98	10,21	6,03
vrpnc5	1,84	90,9	0,57	0,00	2,76	55,41	14,28	14,83
vrpnc11	2,75	22,2	0,00	0,00	0,75	22,43	18,63	6,24
vrpnc12	0,00	16,0	0,00	0,00	0,00	7,21	19,07	4,55
Average	0,68	38,01	0,08	0,00	0,88	25,47	10,52	5,48

Table 7.23: Comparison of SRC-GA and UC-GA and proposed results of three Tabu Search heuristics and the Hybrid Genetic Algorithm. The results above the line are obtained using SRC-GA and the results below the line are from using UC-GA. Algorithms SRC-GA and UC-GA do not compete with the other metaheuristics.

The results in the table above illustrate that SRC-GA and UC-GA are outperformed by the other metaheuristics. Even though the comparison is not perfectly reasonable, the results are quite clear because the difference in the performance is rather large. The HGA-VRP appears to be rather effective on average since its results are quite close to the TS heuristics, which results propose their best performance. Of those algorithms that present the time, SRC-GA and UC-GA are definitely faster with an average time that is about $\frac{1}{5}$ of the time HGA-VRP requires.

Courdeau et al. [4] make an analysis of the three Tabu Search heuristics above, among others, in order to evaluate them in terms of accuracy, speed, simplicity and flexibility. The results are illustrated in table 7.24. An effort has been made to evaluate the algorithms SRC-GA and UC-GA together and HGA-VRP as well, with the help of the results in table 7.23. The results can only be used for comparison, since they are mostly based on personal evaluation. The flexibility of HGA-VRP is left empty because its results are only illustrated for the CMT problems in [2].

	Taburoute	Taillard	Adaptive memory	HGA-VRP	SRC-GA/UC-GA
Accuracy	High	Very high	Very high	High	Low
Speed	Medium	Low	Low	Medium	High
Simplicity	Medium	Medium-low	Medium-low	Medium	High
Flexibility	High	High	High	?	High

Table 7.24: Evaluation of SRC-GA, UC-GA and other metaheuristics.

When it comes to accuracy neither SRC-GA nor UC-GA perform well enough. On the other hand, they score high on speed, simplicity and flexibility.

7.4   Summary

In this chapter the results of the testing have been illustrated. Combination 7 with UC and SIA turned out to be the best one for the fast algorithm when solving the small problems. For the slow algorithm, combination 1 with SRC and SIA gave the best results for the small problems. The slow algorithm outperformed the fast one and was used for further comparison. Combination 7 provided the best results for both the fast and the slow algorithm when solving the large problems. The slow algorithm outperformed the fast one. The best algorithms for the small and the large problems were named SRC-GA and UC-GA and compared to current best metaheuristics. They were not accurate enough, but they were quite fast. In the next chapter the results are discussed.

Chapter 8

Discussion

In this chapter the results of the testing, as illustrated in previous chapter, are discussed. The discussion is divided into five sections. The first four sections discuss the results of the comparison of different combinations for both problem sizes and both types of algorithms. The last one discusses the results in general and the final results.

8.1 Small Problems and Fast Algorithm

When the fast algorithm is applied to the small problems, SRC and UC perform best. BOC and HLC provide a bit worse solutions. The pattern is the same whether the operators are supplied with or without SIA. It is interesting how much effect SIA has. The influence SIA has on HLC and UC indicates that in both operators considerably many customers are inserted into the offspring by the means of the Sweep Algorithm.

It is also interesting how SRC outperforms BOC because BOC was actually thought of as an extended version of SRC. It seems as if BOC performs particularly badly on B-n68-k9 compared to the other two problems. In order to try to find out why, the geography of problems, B-n41-k6, B-n68-k9 and vrpnc2, is first considered. Figures 8.1, 8.2 and 8.3 show the position of the customers and the depot of the problems.

Figure 8.1: The position of the customers and the depot for B-n41-k6. Note that the depot is the diamond in (37,35).

Figure 8.2: The position of the customers and the depot for B-n68-k9. Note that the depot is the diamond in (87,39).

Figure 8.3: The position of the customers and the depot for vrpnc2. Note that the depot is the diamond in (40,40).

The figures illustrate that the customers of B-n41-k6 and B-n68-k9 are gathered in clusters and the customers of vrpnc2 are uniformly distributed. BOC does not appear to have problems with solving B-n41-k6. Thus, one can conclude that the relatively bad performance for B-n68-k9 is not due to the geography of the problem. It is therefore reasonable to compare the optimal solutions of B-n41-k6 and B-n68-k9. Figures 8.4 and 8.5 show the optimal solutions of B-n41-k6 and B-n68-k9.

Figure 8.4: The optimal solution of B-n41-k6.

Figure 8.5: The optimal solution of B-n68-k9.

No visible difference is between the structure of the solutions but when the utilisation of
the capacity is observed the difference becomes apparent. The capacity limit is 100 for
both problems. The total demands of the routes in B-n41-k6 are:

1	2	3	4	5	6
95	100	94	100	99	79

And the total demands of the routes in B-n68-k9 are:

1	2	3	4	5	6	7	8	9
100	100	47	90	99	98	95	100	100

The optimal solutions of both problems have one route that has relatively small total demand, i.e. route 6 in B-n41-k6 and route 3 in B-n68-k9. However, the total demand of route 3 in B-n68-k9 is much smaller, where it uses less than half of the available capacity. The conclusion is that BOC has some difficulties with obtaining good enough solution for problems of the same kind as B-n68-k9 is.

8.2   Small Problems and Slow Algorithm

Combination 1 with SRC and SIA is the best one when using the slow algorithm to solve the small problems. But combinations 3 with BOC and SIA and 7 with UC and SIA also perform well. The effect of SIA is just as clear here as it was for the fast algorithm.

As was explained in the previous section, BOC has some problems with solving problem B-n68-k6. Although, it appears as when BOC is applied with the slow algorithm, the difference in performance of the problems B-n68-k9 and vrpnc2 becomes quite smaller compared to when it is applied with the fast algorithm, see tables 7.14 and 7.19.

8.3   Large Problems and Fast Algorithm

It is interesting how HLC and UC clearly outperform SRC and BOC when the problems become larger. Apparently, the randomness included in both SRC and BOC does not provide accurate enough results and more sophisticated methods are needed to obtain good solutions.

The performance of the four operators without SIA is quite poor. HLC performs particularly bad without SIA. That strongly indicates that too many customers are inserted into the routes of the offspring by the means of the Sweep Algorithm. UC does not seem to rely as much on SIA, thus it can be concluded that relatively few routes are moved directly to the offsprings from the parent solutions.

It is also interesting how in general the algorithm is better able to solve vrpnc4 than vrpnc11, because vrpnc4 has 150 customers but vrpnc11 has only 120 customers. In figures 8.6, 8.7 and 8.8 the structure of the three problems is illustrated.

Figure 8.6: The position of the customers and the depot for vrpnc3. Note that the depot is the diamond in (35,35).

Figure 8.7: The position of the customers and the depot for problem vrpnc4. Note that the depot is the diamond in (35,35).

Figure 8.8: The position of the customers and the depot for problem vrpnc11. Note that the depot is the diamond in (10,45).

In problem vrpnc4 the customers are uniformly distributed but in problem vrpnc11 they are gather in clusters. It appears as the distribution of the customers affects the performance here. The best known solutions of the problems were not found so different structure of the solutions can not be ruled out.

8.4 Large Problems and Slow Algorithm

The slow algorithm with combination 7 performs best for the large problems. Combinations 1, 3 and 5 do also perform relatively well. It is interesting to see how the effect of SIA is reduced compared to the fast algorithm. Combination 8 actually provides quite good solutions without SIA.

As for the fast algorithm, most combinations perform worse on vrpnc11 than vrpnc4, even though vrpnc11 has fewer customers. An effort was made to explain the difference in the previous section.

8.5 The Results in general

In general, combination 7 with UC and SIA is the most effective combination, since it was the best one 3 times out of 4. For the small problems combination 1 performs also quite well but combination 5 is the second best for the large problems. The Local

Search algorithm SIA is essential to obtain an effective algorithm. The performance of the combinations without it is unacceptable. The overall performance of BOC is rather bad.

Unfortunately, neither SRC-GA nor UC-GA are competitive to the best TS heuristics or the HGA-VRP. The average performance of the algorithms together was 10,52% from the optimum or best known values where the other metaheuristics are within 1%. Even though the numbers for the TS heuristics are best of several runs it is estimated that difference in performance of 10% is too much for the algorithms to be considered competitive. Also, the results of HGA-VRP are average numbers of several runs and they are also within 1% from the optimum or best known values.

It is unfortunate that the results of the TS heuristics and HGA-VRP are only presented by the means of the CMT problems. The algorithms SRC-GA and UC-GA is not adjusted particularly to the CMT problems and they were tuned using other types of problems as well. It is interesting to see how UC-GA performs relatively well on problem G-n262-k25, see table 7.22. That indicates that the algorithm does not become less effective as soon as the problems become larger.

In table 7.24 the heuristics are compared by the means of the criteria speed, simplicity and flexibility as well as accuracy. Even though the table is mostly based on personal evaluation, it shows that the algorithms have something to contribute, because it is relatively fast and moreover it is rather simple. Therefore, it would be very interesting to see if some kind of a combination of HGA-VRP and UC with SIA would be promising. The idea of HGA-VRP was shortly introduced in section 2.1.4. It uses so-called parallel version of GA, which generates a number of subpopulations that evolve independently using a serial GA, like has been done here. Once in a while the subpopulations interact. HGA-VRP generates two subpopulations. For instance, in order to allow the algorithm to obtain good solutions for problems of similar type as B-n68-k9, different operators could be used for different populations. Using Geographical Merge on the one population and not on the other will most likely help to obtain good solutions, where one route only uses a small part of the capacity. Therefore, an algorithm would be obtained, which is able to adapt to different problems with different kinds of optimal solutions.

8.6   Summary

In this chapter the results have been discussed. Combination 7 with UC with SIA has provided some promising results, particularly for the larger problems. Although, the results are not quite competitive with proposed results of TS heuristics and the Hybrid Genetic Algorithm. For further work, a combination of HGA-VRP and UC with SIA is suggested. In the next chapter the conclusion in presented.

Chapter 9

Conclusion

The aim of this project was to design an efficient Genetic Algorithm in order to solve the Capacitated Vehicle Routing Problem (VRP). A program was developed based on a smaller program, with rather simple crossover and mutation operators, named Simple Random Crossover and Simple Random Mutation. That program was designed by the author and Hildur Ólafsdóttir in the course Large-Scale Optimisation at DTU in the spring of 2003.

At first, three Local Search Algorithms were designed and implemented, in order to improve single routes in the VRP, one at a time. The algorithms were named Simple Random Algorithm, Non Repeating Algorithm and Steepest Improvement Algorithm. The algorithms were compared based on 10 TSP problems with up to 50 customers. Simple Random Algorithm performed worst by far. The average difference from optimum was $88,40\pm37,78\%$. Non Repeating Algorithm and Steepest Improvement Algorithm provided good results or $5,27\pm2,81\%$ and $4,86\pm2,97\%$ from optimum. Steepest Improvement Algorithm was chosen for further use.

Three new crossover operators were developed; Biggest Overlap Crossover, Horizontal Line Crossover and Uniform Crossover. In different ways, they all focus on the geography of the problem, in order to try to provide good results. In the development process, some drawbacks of the crossover operators were discovered. Therefore, two supporting operators were made. Repairing Operator was designed for Simple Random Crossover and Biggest Overlap Crossover, and for Horizontal Line Crossover and Uniform Crossover Geographical Merge was developed. Eight combinations of operators were defined. Two combinations for each crossover operator, with and without Steepest Improvement Algorithm.

The test problems were divided into small problems defined as having less than 100 customers and large problems with 100 customers or more. Additionally, a fast algorithm with 10000 iterations and relatively small populations and a slow algorithm using 100000 iterations and larger population were defined. For both the small and the large problems, the combinations of operators were compared using both fast and slow algorithm. Three test problems of each size were used. When the fast algorithm was used to solve the small problems, combination 7 with Uniform Crossover and Steepest Improvement Algorithms

performed best. The difference from optimum or best known values was 11,28±3,48%. Combination 1 with Simple Random Crossover and Steepest Improvement provided the best results when the slow algorithm was used to solve the small problems. The results were 4,17±1,49% from the optimum or best know values. For the large problems, combination 7 gave the best results when using the fast algorithm. That resulted in 16,36±2,17% from best known values. When the slow algorithm was used for the large problems, the best results were again obtained by combination 7. The difference from best known values was 10,70±1,21%.

For each size, a choice was made between the fast and the slow algorithm. The algorithms for the small problems were applied to additional 9 problems and the algorithms for the large problems were applied to additional 4 problems. (Thus the testing was made with 12 small problems and 7 large problems.) For the small problems, the fast algorithm was 11,37±3,17% from the optimal or best known values and the slow algorithm was 4,16±1,22%. Furthermore, the fast algorithm was 17,17±2,46% from the optimal or best know values and the slow one was 11,20±1,79%, for the large problems. Thus, the slow algorithm with combination 1 was chosen for the small problems and the slow algorithm with combination 7 was chosen for the large problems. The algorithm for the small problems was called SRC-GA and the algorithm for the large problems was called UC-GA.

The following hypothesis was stated in chapter 1:

It is possible to develop operators for Genetic Algorithms efficient enough to solve large Vehicle Routing Problems.

In order to verify the hypothesis, SRC-GA and UC-GA were compared to the Hybrid Genetic Algorithm (HGA-VRP) and three Tabu Search heuristics: Taburoute, Taillard's Algorithm and Adaptive memory. These heuristics have proposed good results on problems referred to as the Christofides, Mingozzi and Toth problems. The comparison was based on the results of 7 problems. Taburoute, Taillard's Algorithm and Adaptive memory were 0,68, 0,08 and 0,00% from optimum or best known values and HGA-VRP was 0,88%. The results of Taburoute, Taillard's Algorithm and Adaptive memory are the best from several runs but the results of HGA-VRP are average results of several runs. Together SRC-GA and UC-GA did not provide good enough results on these problems, with an average performance of 10,52%. Thus the hypothesis was **rejected**.

However, SRC-GA and UC-GA are on average considerably faster than the other heuristics and more importantly they present some very simple operators. Furthermore, they are rather flexible. For further work focusing on large problems, it could be very interesting to make some other hybrid genetic algorithm with Uniform Crossover and the corresponding operators. That would result in an algorithm with a simple crossover and a number of subpopulations that are each maintained parallel instead of serial, as was done here. Hopefully, it would be able to provide relatively good results more quickly, compared to the one presented in this project.

Bibliography

[1] E. Aarts and J. K. Lenstra. *Local Search in Combinatorial Optimization.* Springer-Verlag, 1996.

[2] J. Berger and M. Barkaoui. A new hybrid genetic algorithm for the capacitated vehicle routing problem. *Journal of the Operational Research Society*, 41(2):179–194, 2003.

[3] Charles Darwin. Britannica concise encyclopedia from encyclopædia britannica., 2004. URL `http://concise.britannica.com/ebc/article?eu=387589`.

[4] J-F. Cordeau, M. Gendreau, G. Laporte, J-Y. Potvin, and F. Semet. A guide to vehicle routing heuristics. *Journal of the Operational Research Society*, 53:512–522, 2002.

[5] Evolutionary Algorithms. Lecture slides: Evolutionary algorithms, 2004. URL `http://www.imm.dtu.dk/courses/02715/`.

[6] E. Falkenauer. A hybrid grouping genetic algorithm for bin packing. *Journal of Heuristics*, 2:5–30, 1996.

[7] M. Fisher. Vehicle routing. *Handbooks of Operations Research and Management Science, chapter 1*, 8:1–31, 1995.

[8] M. Gendreau, A. Hertz, and G. Laporte. New insertion and postoptimization procedures for the traveling salesman problem. *Operations Research*, 40:1086–1094, 1994.

[9] M. Gendreau, A. Hertz, and G. Laporte. A tabu search heuristic for the vehicle routing problem. *Management Science*, 40:1276–1290, 1994.

[10] M. Gendreau, G. Laporte, and J-Y. Potvin. Metaheuristics for the vehicle routing problem. *Technical Report G-98-52, GERAD*, 1999.

[11] M. Jünger, G. Reinelt, and G. Rinaldi. The travelling salesman problem. *Handbooks of Operations Research and Management Science, chapter 4*, 7:225–330, 1995.

[12] G. Laporte, M. Gendreau, J-Y. Potvin, and F. Semet. Classical and modern heuristics for the vehicle routing problem. *International Transactions in Operational Research*, 7:285–300, 2000.

[13] G. Laporte and F. Semet. Classical heuristics for the vehicle routing problem. *Technical Report G-98-54, GERAD*, 1999.

[14] Z. Michalewicz, editor. *Genetic Algorithm + Data Structures = Evolution Programs, 3rd, revised and extended edition.* Springer-Verlag, 1996.

[15] F.B. Pereira, J. Tavares, P. Machado, and E. Costa. GVR: a new genetic representation for the vehicle routing problem. *Prodeedings of the 13th Irish International Conference on Artificial Intelligence and Cognitive Science*, pages 95–102, 2002.

[16] S.M. Sait and H. Youssef, editors. *Iterative Computer Algorithms with Application*

in Engineering: Solving Combinatorial Optimization Problems, chapter 3. IEEE Computer Society, 1999.

[17] The VRP Web. The vrp web, 2004. URL `http://neo.lcc.uma.es/radi-aeb/WebVRP/index.html?/results/BestResults.h%tm`.

[18] Thomas Stidsen. Thomas stidsen, 2003. URL `thomas@stidsen.dk`.

[19] TSPLIB. Tsplib, 2001. URL `http://www.iwr.uni-heidelberg.de/groups/comopt/software/TSPLIB95/`.

[20] Vehicle Routing Data Sets. Vehicle routing data sets, 2003. URL `http://branchandcut.org/VRP/data/`.

[21] L.A. Wolsey, editor. *Integer Programming.* John Wiley and Sons, 1998.

[22] A.H. Wright and J.E. Rowe. Continuous dynamical system models of steady-state genetic algorithms. *Foundations of Genetic Algorithms*, 6:209–226, 2001.

Appendix A

Optimal Values for the Problem Instances in Chapter 3

Problem	Optimal value
kroD100-05	5902
kroD100-10	8553
kroD100-15	9997
kroD100-20	10922
kroD100-25	12173
kroD100-30	13480
kroD100-35	14065
kroD100-40	14542
kroD100-45	14707
kroD100-50	16319

Appendix B

Results of Testing of Repairing Operator in Chapter 4

B.1 Simple Random Crossover

Problem sizes	Difference from optimum			
	Diff. from optimum (%)	Diff. from optimum (%)	Diff. from optimum (%)	Diff. from optimum (%)
32	8,21	23,33	33,49	31,53
44	5.84	17.78	20.5	29.68
60	6.2	18.08	23.63	29.13
80	5.84	14.92	23.59	27.9
Average	6,52	18,53	25,30	29,56

Problem sizes	Standard deviation			
	Std. dev. σ	Std. dev. σ	Std. dev. σ	Std. dev. σ
32	4,45	5,5	6,17	5,01
44	1,38	2,38	3,62	6,16
60	1.69	4,08	4,26	4,41
80	1,74	5,26	7,13	3,98
Average	2,32	4,31	5,30	4,89

B.2   Biggest Crossover Operator

Problem	Difference from optimum			
sizes	Diff. from optimum (%)	Diff. from optimum (%)	Diff. from optimum (%)	Diff. from optimum (%)
32	11,76	22,39	26,99	34,42
44	6,56	18,85	23,5	32,33
60	6,71	20,41	22,3	32,48
80	10,03	18,23	25,92	32,66
Average	8,77	19,97	24,68	32,97

Problem	Standard deviation			
sizes	Std. dev. σ	Std. dev. σ	Std. dev. σ	Std. dev. σ
32	7,95	4,42	4,58	4,73
44	3,05	6,64	3,04	2,12
60	4,29	5,24	5,88	3,74
80	1,85	3,79	3,96	1,95
Average	4,29	5,02	4,37	3,14

Appendix C

Results of Parameter Tuning

C.1 Combination 1, SRC, SRM, RO and SIA

C.1.1 Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	15,48	11,56	9,62	5,28	5,94	1,39	5,66	0,8
50	5,99	1,6	3,14	2,19	5,89	0,51	3,47	2,3
100	4,8	4,08	5,15	2,19	3,42	2,26	3,9	2,64
200	14,52	4,52	8,88	2,31	6,91	1,4	6,48	3

A-n32-k5.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	13,57	3,21	6,76	2,44	5,79	3,39	5,36	2,57
50	8,65	6,76	4,29	1,86	4,21	2,21	3,93	2,26
100	6,51	0,67	5,1	1,79	4,72	2,3	5,84	0,1
200	10,59	3,62	6,35	1,7	6,63	1,34	6,33	2,25

A-n32-k5.vrp, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	13,62	8,85	4,06	2,32	5,77	0,69	5,92	1,06
50	7,14	3,85	4,59	1,78	5,69	2,38	2,7	1,98
100	9,64	3,93	3,9	3	3,39	2,37	3,06	2,34
200	22,32	5,45	9,95	2,32	7,76	5,77	7,53	3,25

vrpcn1, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	44,51	8,07	16,63	3,89	15,3	0,9	7,94	3,12
50	14,98	8,82	11,37	5,39	7,49	4,09	10,54	4,82
100	23,43	13,81	13,46	4,27	10,6	5,46	7,3	5,15
200	18,86	5,52	12,38	3,5	21,71	6,12	12,25	2,61

vrpcn1, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	25,84	23,69	8,44	5,47	8,89	1,86	12,44	6,16
50	11,3	1,68	9,78	3,85	8,95	2,02	6,1	1,79
100	10,41	1,47	14,67	6,06	7,17	2,8	5,97	1,51
200	22,98	2,79	13,84	1,44	18,48	3,16	20,19	5,93

vrpcn1, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	30,22	17,81	11,87	4,63	9,33	4,82	6,16	0,72
50	17,59	4,98	11,37	2,11	12,06	1,82	8,63	3,86
100	18,1	0,87	8,95	3,97	10,48	1,48	14,41	7,35
200	30,16	3,38	20,89	1,4	22,86	2,57	19,94	1,8

C.1.2 Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	12,19	5,66	8,27	1,92	9,85	2,23	4,74	2,06
50	6,38	1,95	3,32	3,22	2,86	2,61	1,79	2,05
100	2,73	1,75	4,62	3,07	2,93	3,05	1,48	2,14
200	3,44	3,56	3,88	2,44	2,91	2,29	1,81	1,92

A-n32-k5.vrp, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	8,62	5,93	7,78	2,63	5,46	2,38	4,11	1,26
50	2,22	1,55	4,13	2,75	4,31	2,07	1,12	2,24
100	2,24	1,1	2,7	1,76	3,32	2,13	0,92	0,75
200	2,6	3	2,45	2,32	1,76	2,3	2,96	2,65

vrpcn1, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	17,33	2,03	6,73	1,45	7,94	2,36	0	0,54
50	2,67	1,65	1,08	2,09	2,67	0,87	1,84	1,21
100	3,49	1,57	3,17	1,48	3,68	2,68	2,16	1,82
200	6,41	0,18	8,32	1,06	3,81	1,12	4,25	1,71

vrpcn1, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	4,57	1,36	5,97	2,7	3,3	3,19	3,68	2,65
50	1,84	0,63	1,21	1,04	0,7	0,5	1,14	1,56
100	5,59	1,47	1,21	0,63	2,92	2,18	2,1	1,27
200	7,43	0,62	7,68	3,27	7,11	0,73	4,57	0,56

C.1.3   Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	17,33	1,95	16,42	4,13	12,16	2,35	12,92	2,3
50	14,3	2,35	16,65	2,55	16,09	1,33	13,39	2,71
100	14,8	3,94	16,24	3,35	13,48	1,87	14,1	2,99
200	23,67	3,83	18,47	3,71	18,24	1,56	18,94	2,27

P-n101-k4.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	16,56	3,55	13,77	2,4	15,62	3,91	12,25	2,16
50	18,18	4,96	15,45	3,77	14,92	2,26	14,19	2,24
100	18,15	1,52	15,62	1,62	18,91	2,66	16,45	2,39
200	23,94	0,84	20,53	2,39	22,64	2,32	20,91	2,59

vrpnc12, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	60,33	5,73	38,86	6,4	26,38	3,07	21,59	9,62
50	64,07	2,82	45,37	2,06	44,67	7,9	32,11	3,63
100	55,24	4,2	38,78	6,42	37,89	5,05	33,74	0,5
200	93,78	4,69	56,34	3,22	66,79	5,38	44,27	3,8

vrpnc12, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	82,93	25,96	41,75	6,11	42,93	5,88	31,67	1,97
50	77,36	14,71	52,97	3,97	51,46	12,39	42,03	2,65
100	76,99	5,5	49,92	2,07	60,69	8,83	47,03	2,21
200	105,69	4,37	76,83	4,47	73,5	5,57	63,01	8,54

C.2   Combination 2, SRC, SRM and RO

C.2.1   Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	21,61	11,68	12,45	3,81	6,56	3,22	11,99	4,64
50	12,53	7,27	9,95	4,4	9,87	7,39	10,89	3,82
100	13,7	4,66	11,94	2,82	9,97	5,1	10,71	5,98
200	24,01	4,2	20,89	1,9	18,29	5,01	18,7	2,12

A-n32-k5.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	17,7	5,65	10,41	2,45	8,78	3,1	10,74	2,47
50	14,13	5,9	9,11	4,86	8,93	4,97	9,44	3,48
100	14,72	5,18	11,4	3,98	7,6	2,99	10	3,25
200	26,58	6,7	20,71	3,96	21,2	2,24	19,54	3,27

A-n32-k5.vrp, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	16,79	8,95	6,91	3,41	8,52	1,43	10,59	1,48
50	15,82	2,43	11,02	2,83	13,37	2,52	12,81	3,4
100	21,96	5,26	10,43	4,47	11,94	3,79	11,63	1,76
200	31,51	3,66	23,98	3,25	23,67	6,5	19,06	6,14

vrpcn1, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	33,02	11,5	17,02	2,89	21,4	3,96	15,68	5,14
50	14,73	1,98	22,92	11,49	18,73	3,9	18,1	2,35
100	24,51	2,82	19,11	5,09	21,02	4,22	16,38	3,53
200	36,57	5,26	31,81	9,83	34,29	3,16	31,56	1,49

vrpcn1, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	37,84	11,12	22,22	7,6	19,49	8,55	19,62	1,48
50	28,25	7,18	23,43	3,19	18,54	6,24	18,67	1,89
100	28,32	3,61	24,44	3,28	20,19	1,89	26,54	5,94
200	42,98	6,67	32,38	2,29	33,52	3,7	38,35	2,36

vrpcn1, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	29,78	9,48	16,06	1,19	21,9	6,95	12,51	3,03
50	37,21	2,11	29,14	2,66	26,67	2,06	28,44	0,65
100	43,11	3,43	26,98	2,36	24,95	4,03	24,25	3,04
200	46,29	4,41	47,68	2,67	39,37	4,34	37,78	5,46

C.2.2   Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	9,9	4,97	8,42	1,33	8,32	4,71	5,48	2,5
50	3,57	2,66	2,98	1,95	3,88	3,26	5,33	3,37
100	5,64	4,21	3,93	2,26	3,6	1,28	3,47	2,75
200	8,7	3,69	7,7	1,36	8,11	2,71	8,62	4,05

A-n32-k5.vrp, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	12,63	3,83	6,94	3,67	6,58	1,41	3,65	2,14
50	4,77	3,04	2,55	3,15	2,83	2,61	3,19	1,88
100	5,48	1,86	3,44	4,6	2,3	0,65	4,64	2,14
200	11,99	4,4	10,31	3,39	7,47	2,44	7,32	2

vrpcn1, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	19,87	10,28	10,48	2,09	10,98	1,45	8,19	2,33
50	9,71	3,06	10,73	2,93	9,21	5,82	8,06	0,73
100	9,4	0,18	8,51	2,53	9,14	3,85	7,56	1,39
200	20,38	3,04	15,56	2,44	14,92	1,09	13,14	2,18

vrpcn1, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	29,27	21,63	13,08	3,86	9,46	3,43	8,38	4,04
50	8,32	0,24	10,35	3,75	7,11	1,56	7,68	0,77
100	10,79	1,6	6,86	0,27	6,1	2,7	6,86	3,91
200	19,56	6,21	16,7	2,09	14,6	2,03	16,7	1,9

C.2.3   Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	34,89	5,84	25,87	5	28,49	7,69	23,32	2,68
50	32,86	7,7	35,83	2,39	33,07	4,88	36,95	4,16
100	31,07	5,73	33,33	3,67	36,04	3,29	33,48	5,11
200	49,07	7,57	40,53	3,42	46,73	5,08	45,58	3,11

P-n101-k4.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	47,49	3,32	44,38	4,68	40,62	4,38	38,65	4,63
50	54,27	6,67	49,07	2,85	49,81	7,55	48,34	7,22
100	55,48	4,69	51,89	6,93	47,72	4,16	53,13	4
200	65,23	6,41	62,44	6,65	67,25	4,09	59,91	3,37

vrpnc12, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	67,28	10,12	39,76	9,65	49,63	11,31	45,33	15,75
50	100,57	3,91	69,19	7,59	68,21	5,36	66,63	3,11
100	92,28	3,87	77,64	12,74	70,2	1,19	68,41	1,66
200	141,06	12,47	91,75	3,69	95,89	6,12	82,44	6,04

vrpnc12, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	110,16	9,68	61,3	5,4	71,5	2,11	45,33	5,79
50	133,01	13,02	86,5	2,75	98,54	10,05	93,9	7,85
100	148,01	8,72	91,02	9,44	89,02	2,75	85,73	7,68
200	168,01	8,03	113,33	7,55	117,93	6,83	102,68	2,45

C.3   Combination 3, BOC, SRM, RO and SIA

C.3.1   Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	35,74	5,76	29,85	7,54	29,06	3,47	25,08	5,77
50	5,1	1,09	5,84	0,22	7,22	4,21	7,02	2,74
100	7,47	2,93	4,67	1,2	3,75	1,81	6,51	4,61
200	10,61	5,19	6,17	0,98	7,19	3,09	7,17	3,12

A-n32-k5.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	35,2	6,91	18,21	2,87	25,08	9,53	20,46	7,13
50	4,29	2,95	4,64	2,17	3,75	2,49	3,9	4,32
100	6,86	1,06	4,59	1,58	4,59	4,16	3,67	1,77
200	11,05	2,63	8,75	3,07	6,76	3,38	5,03	2,89

A-n32-k5.vrp, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	20,82	8	19,16	7,02	21,25	5,3	15,54	5,2
50	5,61	1,86	3,34	2,68	5,15	1,56	3,44	2,26
100	7,22	1,38	1,96	2,13	3,16	2,22	4,29	2,48
200	15,48	1,37	7,55	1,76	7,68	1,91	7,68	1,82

vrpcn1, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	58,22	17,2	40,44	7,91	42,41	2,89	32,51	8,39
50	15,37	7,14	10,67	4,98	8,63	2,03	11,87	4,13
100	5,59	1,75	12,89	4,27	7,37	3,23	13,9	3,1
200	15,37	1,06	21,71	5,04	14,35	1,1	15,17	2,91

vrpcn1, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	46,1	19,53	18,1	4,67	29,71	4,86	29,4	4,82
50	20,95	8,44	6,92	2,21	14,92	5,37	6,35	0,9
100	13,52	3,56	10,67	2,56	6,92	0,32	9,97	1,45
200	18,54	2,94	13,9	1,48	14,86	2,3	12,51	2,46

vrpcn1, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	24,76	1,24	16,25	3,7	15,37	3,7	20,83	6,18
50	13,46	2,02	6,92	2,07	8,83	0,39	12,13	5,43
100	14,48	2,39	11,3	1,09	6,73	2,48	11,05	2,65
200	26,16	3,51	14,86	2,2	13,08	2,53	16,38	2,43

C.3.2   Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	18,19	8,17	18,27	7,29	19,97	5,42	9,77	3,52
50	2,58	2,54	4,29	2,18	2,47	0,75	4,03	1.94
100	5,18	0,63	2,73	2,12	2,37	2,91	2,55	2,56
200	3,98	1,92	3,24	2	3,6	2,68	2,65	2,52

A-n32-k5.vrp, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	11,81	8,56	11,02	3,42	10	2,58	8,24	2,38
50	4,18	1,78	2,93	2,52	2,6	2,54	1,73	2,06
100	4,82	0,93	4,59	1,55	4,03	2,15	3,7	2,06
200	1,53	1,8	1,79	2,04	3,7	1,91	2,86	1,75

vrpcn1, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	51,56	4,9	29,71	8,48	27,49	10,3	21,08	4,25
50	5,65	0,65	3,49	2,34	2,22	0,63	2,41	1,62
100	6,22	0,78	3,49	2,21	3,43	2,57	1,65	1,19
200	5,27	3,3	4,63	0,55	3,56	1,47	5,52	0,54

vrpcn1, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	14,73	5,46	15,87	4,96	17,14	11,59	14,54	3,9
50	3,49	0,24	1,27	0,78	3,94	1,36	2,1	1,79
100	4,19	1,5	6,1	0,71	3,75	0,65	1,65	1,66
200	6,16	0,65	5,14	1,64	4,83	1,65	6,03	0,91

C.3.3   Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	30,07	6,51	34,04	5,94	26,02	4,07	30,04	3,89
50	14,48	2,6	11,45	2,73	15,33	4,05	12,63	3,22
100	15,42	4,21	14,63	2,84	14,48	3,56	11,92	3,26
200	18,47	2,82	16,09	2,42	14,86	0,86	18,09	4,59

P-n101-k4.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	31,16	4,44	29,81	3,79	30,07	4,21	32,92	3,19
50	17,94	1,95	14,6	2,08	14,86	2,07	13,6	1,76
100	15,95	4,62	14,95	3,47	12,75	3,24	13,01	1,22
200	24,55	3,51	20,38	2,66	21,09	1,59	15,89	1,86

vrpnc12, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	79,35	17,26	58,21	8,81	63,54	10,09	55,04	6,6
50	61,83	12,29	41,26	4,48	37,56	3,52	34,59	3,84
100	59,59	2,65	40,12	3,05	41,02	6,4	40,73	2,99
200	86,06	13,29	58,25	4,55	58,41	9,16	47,64	2,07

vrpnc12, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	102,15	26,52	72,52	6,44	60,33	4,46	47,03	9,72
50	68,98	5,03	46,38	1,59	42,76	0,51	52,52	7,27
100	58,13	3,56	44,55	5,93	46,83	3,14	40,61	2,58
200	81,71	7,92	60,24	2,94	64,19	5,72	55,93	12,02

C.4   Combination 4, BOC, SRM and RO

C.4.1   Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	41,79	8,5	44,92	2,73	38,85	6,94	40,51	7,33
50	11,4	3,04	10,41	3,07	10,08	3,66	8,93	2,79
100	10,69	4,41	9,72	2,32	10,48	3,59	9,54	1,65
200	22,5	3,85	17,04	1,82	17,86	2,88	17,19	5,56

A-n32-k5.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	45,94	5,55	33,37	10,83	33,24	7,96	33,67	2,79
50	13,88	1,18	9,18	2,17	12,02	4,1	10,99	2,6
100	14,52	3,49	9,34	3,15	12,73	4,18	9,54	3,01
200	24,21	4,94	21,15	3,97	20	5,18	20,38	4,21

A-n32-k5.vrp, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	30,31	7,7	26,48	8,58	27,04	5,1	27,17	6,58
50	12,76	4,22	13,67	3,87	15,36	2,63	11,61	2,25
100	14,16	3,87	12,5	4,13	13,24	2,41	11,17	3,56
200	32,5	4	21,38	5,06	21,73	4,33	21,2	1,95

vrpcn1, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	54,35	18,03	58,22	4,49	48,63	6,61	56,38	13,9
50	30,22	15,68	15,62	3,9	12,76	3,83	17,78	3,58
100	15,56	4,8	17,65	3,73	17,59	4,82	16,7	1,8
200	34,41	4,82	37,71	1,17	27,68	0,09	28,7	3,46

vrpcn1, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	55,37	11,5	57,27	5,48	54,86	22,27	36,89	3,95
50	23,87	5,84	22,92	7,01	19,56	2,34	19,81	6,07
100	23,24	5,76	27,62	11,87	24,89	2,11	22,41	11,31
200	40,06	5,34	37,52	5,2	33,27	6,59	27,81	4,02

vrpcn1, M=200, IT=10000

Mutation \ Repair	0		50		100		200	
0	42,22	7,32	29,21	3,35	38,86	8,26	28,89	7,03
50	33,71	5,04	26,35	3,16	26,98	3,04	29,52	3,62
100	38,16	2,38	30,54	2,89	31,81	3,11	24	2,03
200	46,73	4,51	40,13	0,32	40,89	4,5	38,6	2,5

C.4.2   Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	29,34	7,56	18,09	6,7	21,71	7,83	27,22	6,17
50	5,08	2,41	3,7	1,83	4,9	2,07	5,89	3,39
100	4,49	2,49	5,18	1,02	3,78	3,06	5,46	1,16
200	6,86	4,35	9,03	2,21	8,67	3,69	7,45	1,4

A-n32-k5.vrp, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	20,82	2,45	15,99	6,93	19,26	5,96	16,94	5,56
50	5,31	1,21	3,49	2,18	3,21	2,09	5,1	4,55
100	4,67	3,22	4,11	2,04	2,7	1,96	5	4,68
200	11,28	2,05	8,32	3,51	8,39	2,21	6,96	1,5

vrpcn1, M=200, IT=100000

Mutation \ Repair	0		50		100		200	
0	42,6	3,3	34,6	5,82	30,22	5,77	30,03	5,99
50	11,75	1,82	12,32	1,26	9,4	3,16	3,3	2,85
100	12,38	4,84	6,6	2,38	8,38	2,56	6,98	0,7
200	18,29	2,18	16,19	1,95	14,1	1,68	14,22	0,59

vrpcn1, M=400, IT=100000

Mutation \ Repair	0		50		100		200	
0	27,05	10,35	29,46	15,18	32,76	4,85	28,89	3,06
50	10,1	0,41	7,05	3,56	5,97	1,06	7,05	2,24
100	8,95	2,56	9,33	2,57	5,02	2,03	4,89	2,05
200	19,68	2,01	14,67	1,23	16,44	1,71	14,35	3,14

C.4.3   Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	56,65	4,56	50,63	5,89	55,8	8,58	55,01	4,88
50	36,71	5,27	25,7	4,92	29,49	2,67	33,86	4,67
100	30,95	4,48	28,11	1,14	25,35	2,83	29,84	3,96
200	46,75	6,7	40,09	1,56	39,24	1,68	38,33	1,04

P-n101-k4.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	62,47	6,54	62,53	14,32	57,97	6,57	61,7	6,52
50	47,96	5,77	43,64	7,51	40,7	3,18	39,09	4,28
100	46,55	7,67	46,46	5,09	43,82	3,42	41,41	3,23
200	61,85	7,29	52,28	4,44	58,09	4	53,86	5,02

vrpnc12, M=50, IT=10000

Mutation \ Repair	0		50		100		200	
0	138,21	13,89	110,41	20,52	92,68	20,66	79,02	9,1
50	71,54	10,12	67,93	3,62	76,3	9,27	64,15	7,31
100	77,4	2,72	59,31	1,59	62,64	4,95	65,65	0,97
200	116,3	12,31	99,76	1,99	95,61	2,04	93,21	9,05

vrpnc12, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	142,76	23,66	92,56	6,14	78,78	10,85	84,76	2,33
50	105,57	7,3	81,06	1,06	84,15	2,8	74,55	4,41
100	110,93	10,73	86,63	4,37	89,02	9,91	76,99	8,42
200	125,81	7,13	106,99	3,16	110,57	4,62	94,76	8,56

C.5   Combination 5, HLC, SRM, GM and SIA

C.5.1   Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation \ Merge	0		50		100		200	
0	90,66	49,95	218,07	207,9	308,86	364,67	76,95	58,94
50	19,46	3,18	14,46	2,47	18,7	1,36	17,45	2,47
100	20,2	6,67	14,72	3,78	16,61	2,75	15,23	1,66
200	28,29	4,22	27,76	2,79	29,34	3,54	28,44	2,95

A-n32-k5.vrp, M=100, IT=10000

Mutation \ Repair	0		50		100		200	
0	54,15	16,58	101,9	64,17	55,7	15,66	41,05	9,74
50	22,45	3,3	22,81	1,18	22,81	0,71	21,17	3,9
100	22,35	6,43	21,53	4,51	21,53	2,1	21,99	4,63
200	30,59	1,25	31,35	1,84	29,59	5,45	27,63	2,22

A-n32-k5.vrp, M=200, IT=10000

Mutation \ Merge	0		50		100		200	
0	39,21	6,82	43,95	6,21	47,1	14,1	43,74	6,91
50	28,04	2,78	25,28	2,97	25,97	4,34	21,76	3
100	26,28	2,78	27,14	2,32	22,47	1,85	27,02	2,03
200	31,84	4,65	27,6	6,37	26,43	3,48	29,54	1,66

vrpcn1, M=50, IT=10000

Mut. \ Merge	0		50		100		200	
0	1361,92	636,97	2096,19	2694,95	5129,62	6629,39	2060,16	1382,82
50	13,9	1,17	8,06	2,82	6,73	1,1	7,87	3,47
100	7,75	1,57	12,32	2,41	12,83	3,76	11,62	1,62
200	13,71	2,44	16,89	3,82	13,71	5,41	17,65	5,79

vrpcn1, M=100, IT=10000

Mutation \ Merge	0		50		100		200	
0	874,5	431,89	825,96	901,68	223,28	62,6	431,77	329,87
50	14,48	0,54	7,81	3,25	11,17	1,71	7,87	2,38
100	11,3	2,53	11,37	4,63	8,63	2,83	8,38	2,44
200	16,63	4,69	13,59	4,09	16,25	2,89	14,67	3,97

vrpcn1, M=200, IT=10000

Mutation \ Merge	0		50		100		200	
0	156,81	39,58	599,28	576,89	206,1	5,44	185,37	11,51
50	10,54	1,56	12	0,62	9,27	3,76	8,13	6,24
100	10,29	5,13	9,21	3,03	10,41	2,09	9,46	3,7
200	19,68	3,76	15,62	3,26	19,37	0,8	15,05	4,61

C.5.2 Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation \ Merge	0		50		100		200	
0	39,67	7,57	172,92	163,37	42,55	5,45	43,49	4,66
50	11,12	2,87	7,93	1,55	10,23	1,79	9,03	2,18
100	10,36	1,95	8,75	1,17	9,97	1,57	8,47	2,79
200	21,45	3,35	24,85	1,29	21,33	2,09	23,14	1,41

A-n32-k5.vrp, M=400, IT=100000

Mutation \ Merge	0		50		100		200	
0	33,7	8,77	37,22	2,45	38,62	5,65	38,7	4,25
50	13,11	1,52	12,24	1,63	11,15	3,47	12,37	1,32
100	13,8	2,07	11,45	1,71	11,07	1,14	10,89	1,69
200	26,68	1,46	22,58	2,7	22,73	2,19	25,33	1,07

vrpcn1, M=200, IT=100000

Mutation \ Merge	0		50		100		200	
0	3528,19	1330,88	2755,74	2682,38	664,98	675,82	1133,81	652,56
50	7,11	3,5	6,16	2,05	6,03	3,23	9,71	2,35
100	8,83	4,46	6,6	2,78	9,46	2,03	8,44	1,04
200	12,95	2,29	9,02	2,38	9,21	2,96	7,05	0,87

vrpcn1, M=400, IT=100000

Mutation \ Merge	0		50		100		200	
0	1139,03	662,73	639,52	628,62	183,49	1,13	184,26	5,64
50	5,52	1,62	7,11	0,36	6,29	0,71	4,57	1,23
100	6,67	3,39	3,94	0,88	6,41	3,46	2,98	1,09
200	10,92	3,44	6,22	2,8	6,67	0,78	8,13	1,97

C.5.3   Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation \ Merge	0		50		100		200	
0	359,76	7,07	367,04	5,63	363,02	6,18	364,97	2,02
50	4,88	1,27	2,97	0,68	3,47	1,56	3,94	1,13
100	4,14	1,34	2,7	0,22	3,32	0,71	2,79	0,54
200	8,37	1,33	8,6	0,73	7,2	1,03	8,11	1,2

P-n101-k4.vrp, M=100, IT=10000

Mutation \ Merge	0		50		100		200	
0	354,25	8,6	359,91	11,65	357,64	4,11	357,35	6,53
50	4,55	0,78	3,58	0,85	3,32	0,38	3,61	0,67
100	5,32	1,17	3,94	0,76	3,38	0,56	3,2	0,3
200	7,87	0,68	7,96	0,63	7,93	0,81	7,31	0,73

vrpnc12, M=50, IT=10000

Mut. \ Merge	0		50		100		200	
0	939,43	621,54	1861,99	9,76	1265,73	875,78	480,49	8,99
50	25,04	2,5	22,97	2,28	23,54	1,9	24,51	0,34
100	27,4	2,29	23,17	2,18	22,07	2,15	24,11	0,47
200	32,52	1,36	30,69	1,31	29,59	1,4	25,65	3,26

vrpnc12, M=100, IT=10000

Mutation \ Merge	0		50		100		200	
0	180,33	215,19	325,41	212,09	181,67	219,06	183,78	215,51
50	26,63	1,04	24,11	2,53	22,72	2,21	21,22	0,8
100	25,24	1,37	22,6	2,05	24,51	1,91	24,35	1,52
200	33,25	0,5	30,45	1,34	29,92	1,13	28,86	0,64

C.6   Combination 6, HLC, SRM and GM

C.6.1   Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation \ Merge	0		50		100		200	
0	286,99	109,2	394,89	480,67	196,12	104,65	205,32	67,9
50	30,26	3	30,43	1,41	31,02	3,25	29,59	1,99
100	29,26	4,59	28,89	2,96	31,66	1,66	33,34	2,83
200	42,32	5,87	44,08	6,39	39,97	4,96	43,8	2,65

A-n32-k5.vrp, M=100, IT=10000

Mutation \ Merge	0		50		100		200	
0	131,15	13,76	99,64	9,13	106,54	17,21	129,94	53,61
50	38,16	1,75	36,38	2,62	37,78	4,22	36,28	2,2
100	37,22	2,49	37,35	2,23	34,62	5,85	34,82	4,39
200	48,06	4,17	46,99	2,84	46,61	2,3	44,82	4,25

A-n32-k5.vrp, M=200, IT=10000

Mutation \ Merge	0		50		100		200	
0	90,57	3,36	91,29	13,4	99,79	17,04	93,04	20,33
50	45,23	4,55	44,52	4,09	50,13	1,63	41,43	4,31
100	43,01	3,15	46,25	3,53	48,37	3,28	40,64	2,42
200	49,95	4,41	49,67	5,01	49,21	4,74	52,14	1,1

vrpcn1, M=50, IT=10000

Mut. \ Merge	0		50		100		200	
0	188,29	212,98	688,59	395,24	260,54	166,85	239,83	191,34
50	7,75	5,07	7,81	2,76	3,05	2,19	8,7	3,2
100	10,22	4,72	6,73	2,99	12,83	4,13	10,29	7
200	13,4	3,09	9,84	2,38	8,38	1,21	14,86	0,82

vrpcn1, M=100, IT=10000

Mutation \ Merge	0		50		100		200	
0	38,04	13,69	48,03	39,06	39,56	22,86	37,17	6,75
50	8,19	1,95	5,46	2,33	10,73	2,07	9,33	3,9
100	8,38	3,56	3,05	2,73	4,13	2,44	6,67	1,02
200	11,94	2,82	9,9	3,42	10,35	4,89	9,71	2,35

vrpcn1, M=200, IT=10000

Mutation \ Merge	0		50		100		200	
0	41,87	8,09	18,1	3,7	29,93	2,48	31	5,79
50	4,63	2,89	8,19	3,82	3,87	3,77	5,46	3,82
100	8,7	4,66	9,14	1,73	5,4	4,57	7,62	2,3
200	12,89	6,06	11,43	4,75	8,89	1,63	14,54	0,55

C.6.2 Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation \ Merge	0		50		100		200	
0	97,6	9,81	140,24	138,33	155,93	129,32	158,38	140,13
50	17,12	2,57	19,54	3,88	18,93	0,9	20,89	2,52
100	20,33	2,74	20,26	3,41	21,2	1,44	19,44	1,85
200	35,99	3,47	36,33	3,23	37,83	0,87	35,15	4,37

A-n32-k5.vrp, M=400, IT=100000

Mutation \ Merge	0		50		100		200	
0	87,96	11,63	77,22	6,19	83,29	6,91	82,27	10,45
50	24,9	3,24	24,74	2,83	25,18	1,77	22,3	3,76
100	25,84	2,72	23,72	2,83	21,73	3,34	25,05	2,5
200	37,81	4,1	39,08	5,55	40,26	2,07	40	2,59

vrpcn1, M=200, IT=100000

Mutation \ Merge	0		50		100		200	
0	1669,87	7,81	2261,78	2903,68	1074,54	1247,3	4329,62	2929,92
50	17,02	4,7	23,81	5,88	19,17	3,09	16,83	3,24
100	16,57	3,21	15,62	3,44	21,59	1,26	21,97	8,02
200	32,63	1,66	32,95	0,97	32,83	4,66	34,41	2,59

vrpcn1, M=400, IT=100000

Mutation \ Merge	0		50		100		200	
0	856,94	586,58	687,9	705,9	185,4	7,41	185,4	5,76
50	22,1	3,12	21,08	1,8	21,27	1,53	24,13	1,86
100	23,81	5,2	23,24	4,36	28,32	3,04	22,48	0,81
200	37,52	2,75	36,76	2,06	36,32	1,32	34,67	3,06

C.6.3   Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation \ Merge	0		50		100		200	
0	450,86	2,87	455,07	7,48	449,31	6,05	449,44	5,57
50	51,34	6,57	39,68	8,93	35,3	5,14	38,68	8,54
100	35,95	3,36	36,33	4,55	33,89	1,2	32,89	6,76
200	63,08	5,15	66,46	8,43	62,94	3,72	65,32	7,89

P-n101-k4.vrp, M=100, IT=10000

Mutation \ Merge	0		50		100		200	
0	447,13	5,09	450,91	4,85	447,89	8,97	446,81	3,75
50	46,08	3,88	35,42	2,16	37,3	3,78	42,79	5,34
100	51,04	9,43	38,88	2,52	38,09	1,69	40,88	3,57
200	62,64	7,58	67,28	1,71	66,7	5,94	61,44	4,33

vrpnc12, M=50, IT=10000

Mut. \ Merge	0		50		100		200	
0	811,3	753,86	966,79	632,5	818,13	776,42	1286,83	873,05
50	48,74	1,26	39,07	5,59	40,98	4,15	44,92	2,91
100	49,51	2,37	44,11	1,1	46,22	1,83	43,29	3,74
200	58,29	2,48	50,2	3,34	56,54	4,35	54,76	2,28

vrpnc12, M=100, IT=10000

Mutation \ Merge	0		50		100		200	
0	214,51	216,89	355,81	208,1	198,98	209,92	51,3	1,4
50	49,67	2,95	48,46	3,02	45,28	2,63	45,61	1,25
100	50,28	0,66	47,89	1,84	48,09	1,2	43,66	2,71
200	58,5	2,47	57,52	2,45	52,89	1,79	50,57	0,93

C.7   Combination 7, UC, SRM, GM and SIA

C.7.1   Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation\Merge	0		50		100		200	
0	196,58	131,62	269,83	225,37	306,02	228,42	231,49	117,76
50	9,52	2,86	4,31	2,44	2,98	3,23	6,07	4,17
100	11,35	4,29	5,43	3,81	4,87	2,67	7,65	1,98
200	7,32	3,41	7,98	2,94	6,05	2,04	4,95	2,84

A-n32-k5.vrp, M=100, IT=10000

Mutation\Merge	0		50		100		200	
0	88,39	50,83	83,58	57,37	66,31	18,21	84,34	64,3
50	3,57	2,32	2,98	3,59	1,73	2,31	2,45	2,7
100	6,86	5,08	1,96	2,5	2,35	2,82	3,67	3,23
200	11,38	4,57	7,88	0,86	6,79	2,3	8,19	1,8

A-n32-k5.vrp, M=200, IT=10000

Mutation\Merge	0		50		100		200	
0	65,28	18,45	71,43	7,92	68,95	49,39	33,04	18,89
50	6,2	3,23	4,87	2,84	2,14	2,62	5,79	4,67
100	6,51	1,53	5,23	2,75	3,85	3,11	2,83	2,39
200	10,03	4,31	7,09	3,91	4,34	2,81	6,45	2,49

vrpcn1, M=50, IT=10000

Mutation\Merge	0		50		100		200	
0	217,21	195,86	112,13	66,77	305,81	235,59	450,38	196,97
50	7,3	4,61	12	2,24	9,71	2,24	8,32	8,14
100	9,4	3,14	8,83	6,8	7,05	3,56	9,71	2,7
200	12,19	2,16	13,59	2,38	11,68	4,4	13,52	2,59

vrpcn1, M=100, IT=10000

Mutation\Merge	0		50		100		200	
0	95,04	70,68	102,44	66,02	56,72	30,86	78,44	79,36
50	7,11	1,63	4,44	2,61	8	2,66	7,37	0,45
100	11,43	3,27	4,89	2,68	7,81	2,64	9,71	6,84
200	10,16	5,84	5,46	5,38	12	3,77	8,19	2,18

vrpcn1, M=200, IT=10000

Mutation\Merge	0		50		100		200	
0	37,4	6,3	32,1	6,49	25,59	10,74	31,78	8,8
50	4,83	2,46	5,9	3,93	7,43	1,94	4,19	4,21
100	13,21	3,68	5,33	3,37	3,56	0,91	3,75	4,36
200	15,37	7,21	14,41	4,5	10,41	4,38	11,68	6,02

C.7.2 Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation\Merge	0		50		100		200	
0	174,15	177,24	90,39	145,8	230,08	175,28	156,91	167,29
50	4,85	2,82	1,1	2,19	2,45	3,01	1,91	2,42
100	4,26	2,16	0	0	4,77	0,94	2,3	1,87
200	4,46	2,33	1,94	2,4	0,77	1,53	1,91	2,42

A-n32-k5.vrp, M=400, IT=100000

Mutation\Merge	0		50		100		200	
0	16,58	3,28	14,03	4,2	13,88	2,28	14,95	3,11
50	3,44	2,76	0,03	0,05	2,19	2,21	3,47	2,84
100	2,78	1,7	2,63	2,23	3,01	2,54	0,77	1,53
100	2,58	2,35	0,71	1,43	2,3	2,4	1,02	1,91

vrpcn1, M=200, IT=100000

Mutation\Merge	0		50		100		200	
0	25,65	3,77	21,46	2,75	21,18	1,49	23,98	3,32
50	5,27	0,24	2,29	1,89	1,97	1,36	1,65	2,33
100	3,17	4,22	3,75	3,24	3,24	1,5	1,78	1,1
200	1,84	1,45	3,81	2,18	0,83	1,04	2,98	2,34

vrpcn1, M=400, IT=100000

Mutation\Merge	0		50		100		200	
0	26,1	1,65	16,32	4,2	16,25	3,6	74,02	80,37
50	2,67	1,36	-0,32	0,63	-0,51	0,36	-0,25	0,71
100	4,06	2,12	0,44	1,21	1,84	0,86	0	0,62
200	5,27	1,71	2,1	2,57	1,33	0,68	1,84	1,8

C.7.3　Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation\Merge	0		50		100		200	
0	74,57	3,18	74,6	1,03	72,1	2,38	71,48	3,69
50	10,43	3,35	8,49	2,23	9,72	1,42	7,99	4
100	13,3	1,78	8,4	2,8	8,55	1,06	7,81	3,53
200	15,15	1,71	12,48	3,11	11,89	3,5	9,81	1,77

P-n101-k4.vrp, M=100, IT=10000

Mutation\Merge	0		50		100		200	
0	69,81	2,57	132,7	87,22	70,46	2,38	70,63	3,93
50	12,51	2,1	11,37	2,4	10,19	2,09	11,63	1,87
100	11,25	1,2	7,64	3,07	11,78	1,52	8,81	2,14
200	12,28	2,08	11,07	3,34	12,86	3,83	10,4	3,14

vrpnc12, M=50, IT=10000

Mutation\Merge	0		50		100		200	
0	33,98	4,8	29,72	1,93	176,18	208,97	34,15	6,39
50	29,31	4,07	28,05	2,85	20,93	1,23	25,24	0,72
100	25,04	4,29	23,33	2,5	23,13	2,13	24,72	1,91
200	33,66	1,49	26,26	2,13	27,4	1,06	28,29	2,33

vrpnc12, M=100, IT=10000

Mutation\Merge	0		50		100		200	
0	29,43	2,56	28,9	2,41	28,66	1,35	29,84	2,88
50	28,41	2,33	21,91	4,37	24,55	0,35	23,74	2,74
100	24,76	2,35	24,47	2,67	24,59	2,7	23,37	2
200	31,99	1,84	25,33	1,45	26,67	1,07	30,37	2,8

C.8　Combination 8, UFC, SRM and GM

C.8.1　Small and Fast

A-n32-k5.vrp, M=50, IT=10000

Mutation\ Merge	0		50		100		200	
0	321,15	251,46	182,14	121,11	353,58	481,9	448,54	303,37
50	12,6	4,43	10,28	6,98	6,94	2,92	8,93	4,67
100	15	5,33	12,86	3,91	10,66	4,27	5,71	2,88
200	14,92	2,23	11,07	4,77	13,06	2,73	13,52	3,53

A-n32-k5.vrp, M=100, IT=10000

Mutation\ Merge	0		50		100		200	
0	92,13	47,42	80,17	52,34	116,55	63,02	67,8	11,5
50	10,79	2,6	8,75	3,2	9,57	4,21	6,51	4,5
100	14,52	7,6	8,49	5,68	10,23	2,08	9,18	1,8
200	18,34	1,55	11,07	3,87	15,43	4,45	14,16	2,64

A-n32-k5.vrp, M=200, IT=10000

Mutation\ Merge	0		50		100		200	
0	58,19	19,82	50,47	13,9	64,44	8,73	59,39	8,29
50	9,08	3,41	9,74	5,35	10,74	6,85	11,43	4,62
100	12,24	1,59	5,89	3,19	6,66	2,7	9,41	2,13
200	15,56	3,16	15,36	4,47	16,02	4,8	19,23	6,28

vrpcn1, M=50, IT=10000

Mutation\ Merge	0		50		100		200	
0	225,83	106,32	1118,17	565,71	598,66	442,08	286,81	296,9
50	22,67	0,82	23,11	8,88	18,67	3,77	19,68	3,55
100	19,3	2,99	15,87	0,59	17,08	3,41	21,97	5,26
200	19,75	3,2	18,22	2,11	21,08	1,66	25,65	1,39

vrpcn1, M=100, IT=10000

Mutation\ Merge	0		50		100		200	
0	113,51	26,68	216,41	198,92	66,84	6,24	66,56	8,15
50	19,81	2,85	15,87	4	16,25	3,35	18,1	2,75
100	14,6	2,88	17,52	5,16	15,56	4,98	14,41	6,18
200	25,46	5,11	22,86	2,42	21,27	0,86	27,94	0,48

vrpcn1, M=200, IT=10000

Mutation\ Merge	0		50		100		200	
0	63,06	13,69	56,06	2,56	59,35	2,01	62,18	4,68
50	21,27	4,42	14,41	5,44	19,17	1,32	18,1	2,75
100	23,87	3,73	13,84	1,88	14,98	1,3	18,48	2,75
200	29,52	6,43	29,9	1,94	28,25	3,95	30,29	3,66

C.8.2   Small and Slow

A-n32-k5.vrp, M=200, IT=100000

Mutation\ Merge	0		50		100		200	
0	253,71	161,43	312,18	140,47	179,65	173,59	180,31	174,09
50	9,39	1,94	3,67	1,95	2,98	2,48	3,01	2,72
100	6,66	1,8	4,44	2,67	5,03	1,81	2,42	2,16
200	4,87	1,88	5,66	3,56	5,61	3,32	3,98	1,18

A-n32-k5.vrp, M=400, IT=100000

Mutation\ Merge	0		50		100		200	
0	103,55	131,92	99,34	134,15	26,43	5,36	163,19	164,38
50	8,14	3,49	2,09	2,09	4,13	2,48	4,54	2,14
100	2,65	2,03	5,77	3,16	2,35	2,53	3,11	1,91
200	4,44	2,64	5,92	3,54	4,34	3,29	4,67	2,21

vrpcn1, M=200, IT=100000

Mutation\ Merge	0		50		100		200	
0	59,43	2,42	161,7	83,34	161,38	73,61	343	308,88
50	11,94	0,8	6,54	1,15	9,46	3,34	8,06	4,68
100	7,11	1,8	8,95	4,75	6,48	1,48	7,62	2,7
200	12,83	2,38	14,67	7,15	12,38	5,35	11,43	2,39

vrpcn1, M=400, IT=100000

Mutation\ Merge	0		50		100		200	
0	46,29	1,95	100,32	83,03	152,93	80,96	98,46	85,1
50	12,38	1,27	5,59	2,8	5,14	1,48	6,92	4,38
100	7,81	1,62	6,41	3,59	7,43	3,64	9,9	2,97
200	12,38	0,82	11,17	1,59	11,87	2,6	11,94	1,26

C.8.3   Large and Fast

P-n101-k4.vrp, M=50, IT=10000

Mutation\ Merge	0		50		100		200	
0	380,48	40,92	278,24	110,21	294,39	88,65	368,25	9,69
50	30,95	2,94	34,01	7,66	26,84	5,25	30,34	3,86
100	34,3	5,47	27,72	4,78	29,19	2,67	23,52	5,28
200	36,77	4,17	38,03	5,02	38,68	6,19	34,36	2,49

P-n101-k4.vrp, M=100, IT=10000

Mutation\ Merge	0		50		100		200	
0	353,16	4,2	286,49	88,46	264,85	115,61	1013,66	1388,97
50	38,09	5,04	32,13	2,78	32,13	4,57	30,1	3,86
100	40,18	5,9	35,92	6,86	34,3	5,84	29,37	4,72
200	38,77	4,65	40,56	4,8	38,77	2,77	37,18	3,31

vrpnc12, M=50, IT=10000

Mut.\ Merge	0		50		100		200	
0	48,82	3,91	52,72	2,19	207,4	216,92	46,79	3,61
50	55,28	3,61	46,38	2,73	45,12	4,5	50,53	1,36
100	54,72	3,44	44,07	3,92	37,6	1,36	43,86	5,21
200	58,29	3,98	48,86	5,29	46,99	4,43	50,57	6,24

vrpnc12, M=100, IT=10000

Mutation\ Merge	0		50		100		200	
0	47,97	1,79	40,93	2,91	46,83	2,45	44,35	1,88
50	44,07	2,09	40,28	3,18	37,07	3,88	40,69	3,88
100	49,35	2,88	41,46	3,39	39,19	1,15	42,76	0,96
200	48,74	5,35	44,92	0,98	43,33	2,69	43,9	2,8